



Shadow: Real Applications, Simulated Networks

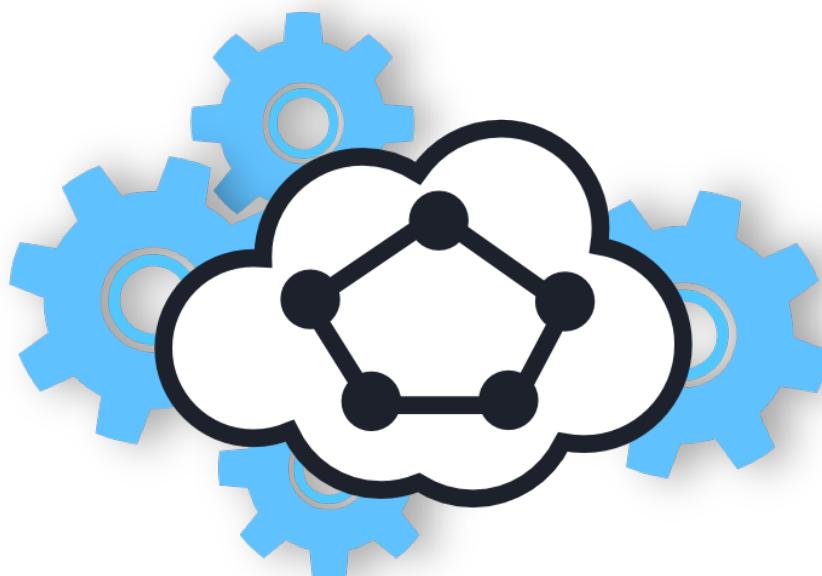
Dr. Rob Jansen

U.S. Naval Research Laboratory
Center for High Assurance Computer Systems

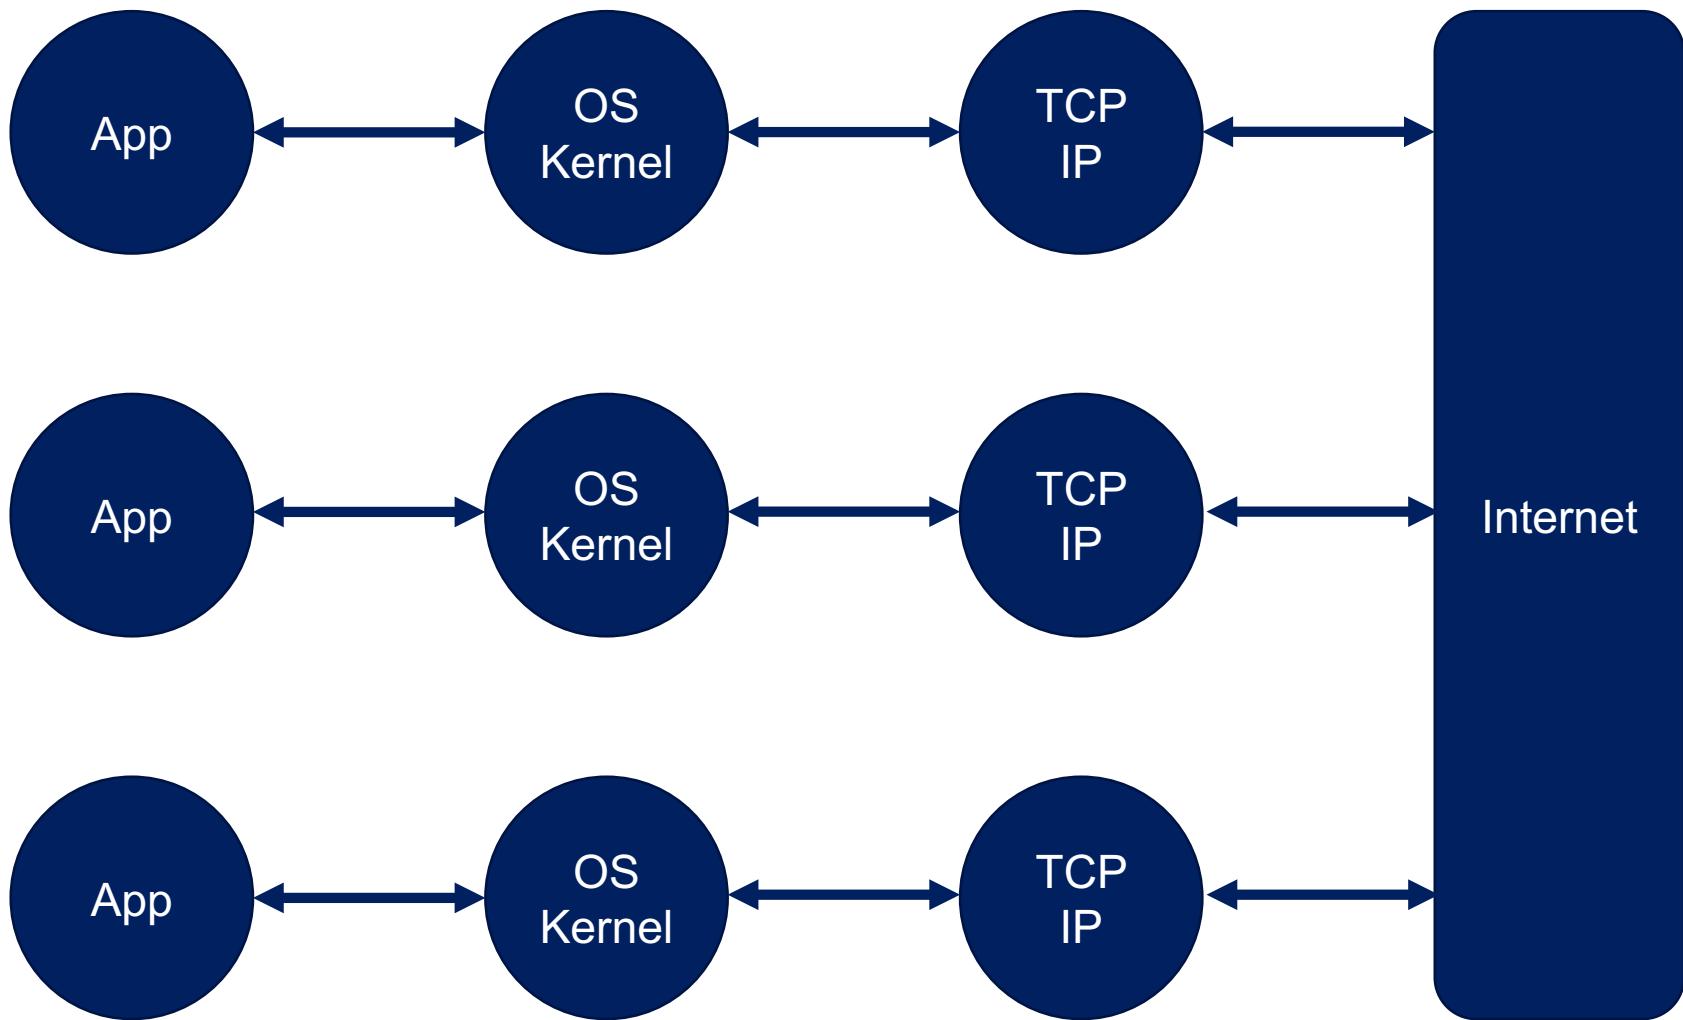
Cyber Modeling and Simulation Technical Working Group
Mark Center, Alexandria, VA
October 25th, 2017

What is Shadow?

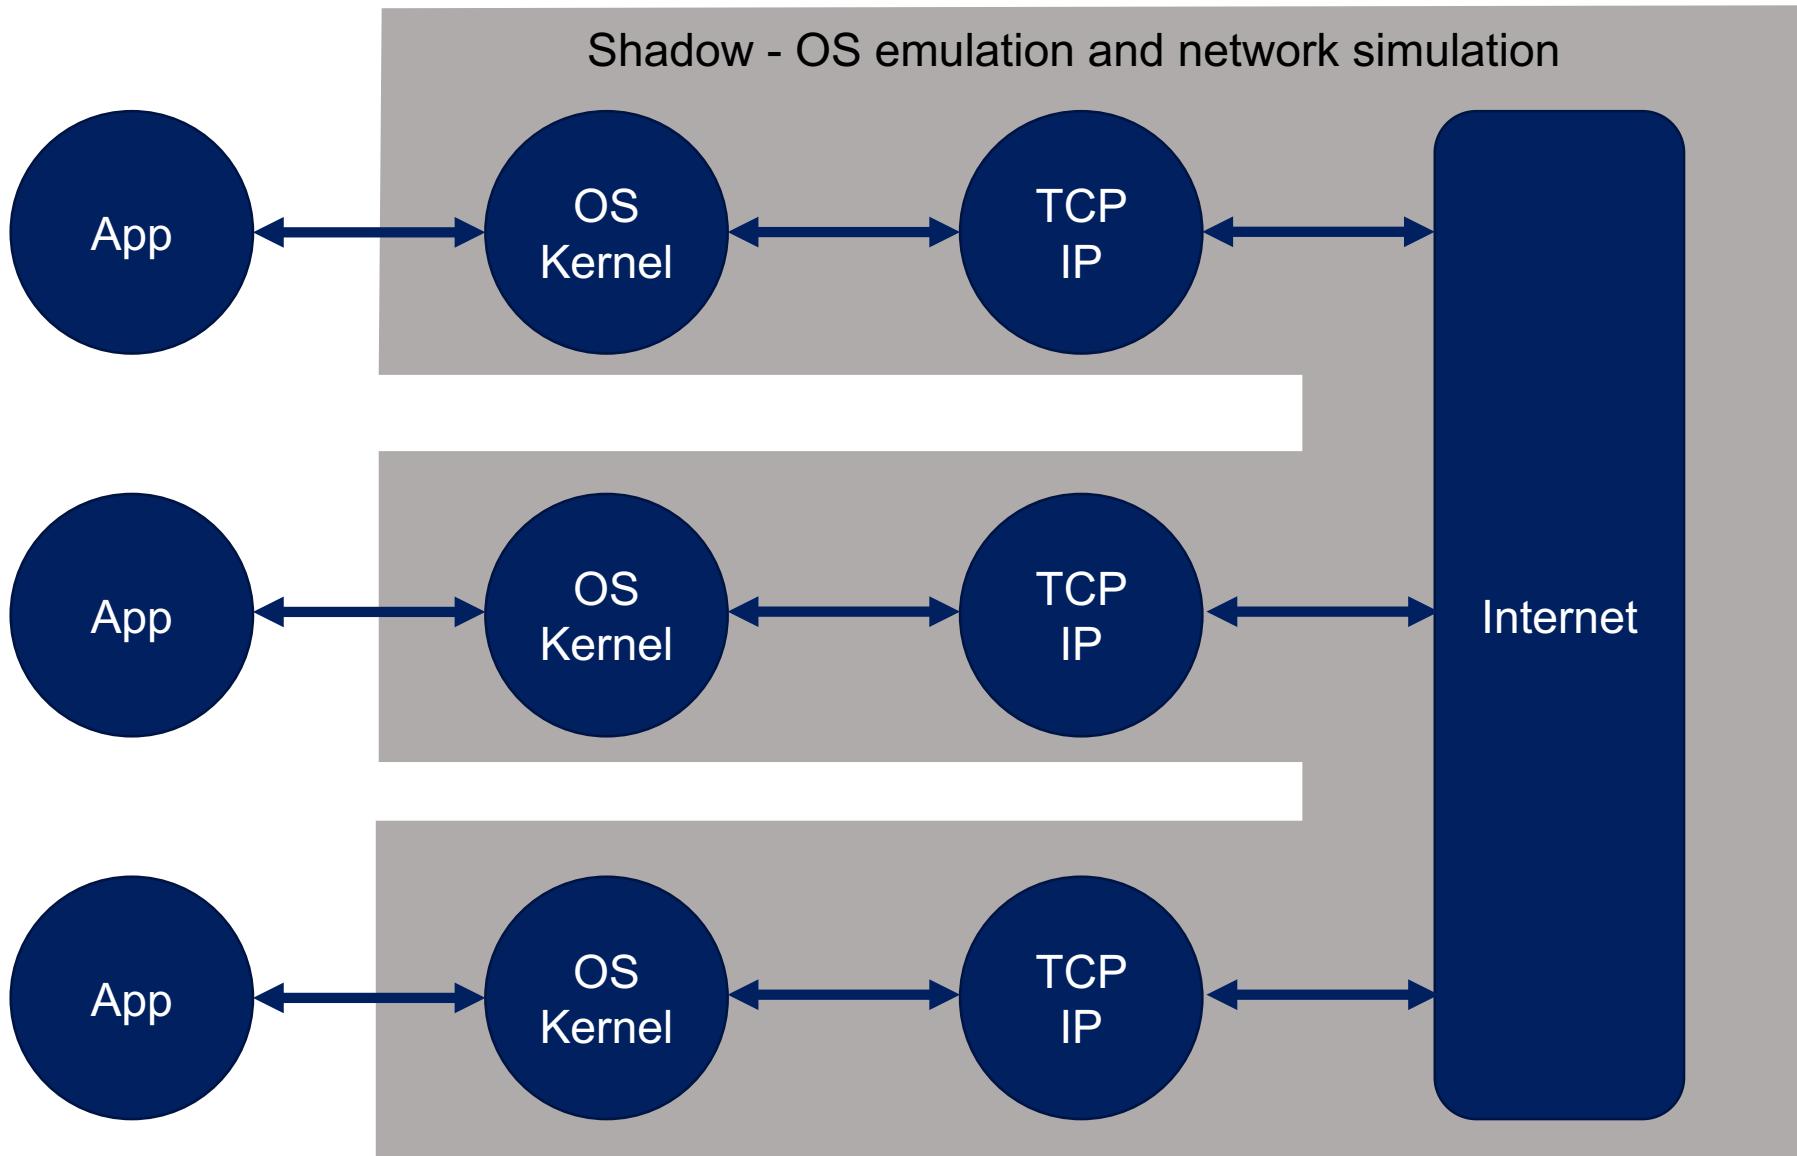
- Open-source network simulator / emulator hybrid
- Directly executes real applications like Tor and Bitcoin over a simulated network topology
- Efficient, scalable, deterministic, accurate, and more!



How does Shadow Work?



How does Shadow Work?



Why should you care?

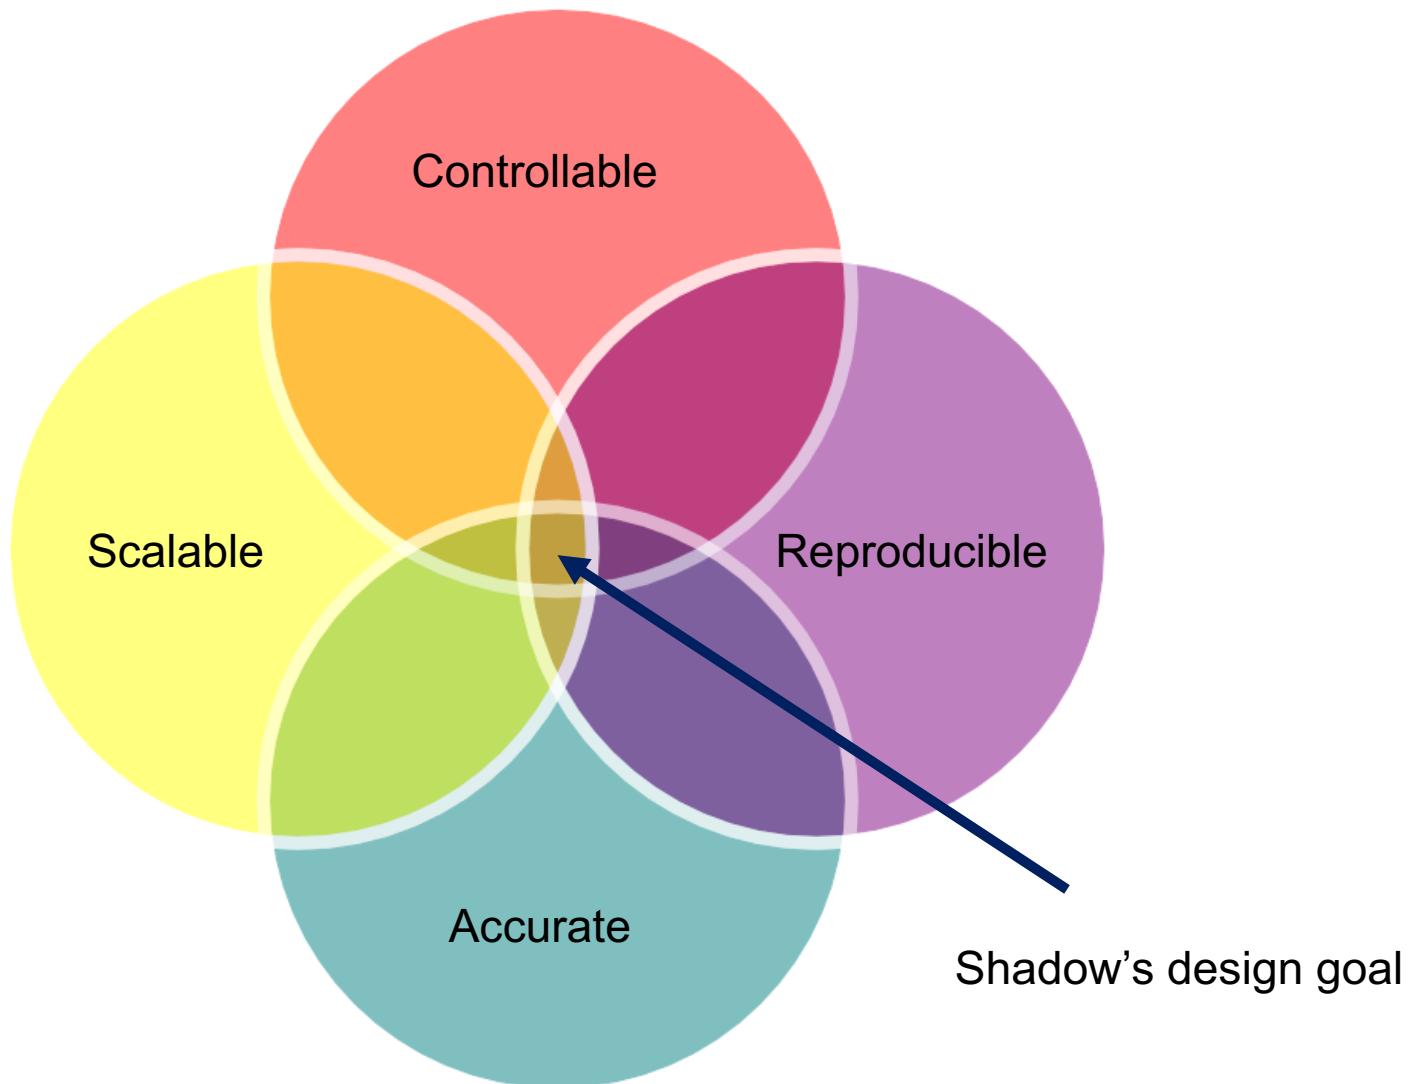
- Expedite research and development
- Evaluate software mods or attacks without harming real users
- Understand holistic effects before deployment
- Shadow supports simulation for general-purpose applications

Talk Outline

- Experimentation options and tradeoffs
- Shadow design
- Simulation use cases – The Tor Anonymity Network
 - Gentle Tor introduction
 - Congestion problems
 - Scheduling algorithms
 - Performance enhancing algorithms
 - Denial of service attacks
- Conclusion

Experimentation Options

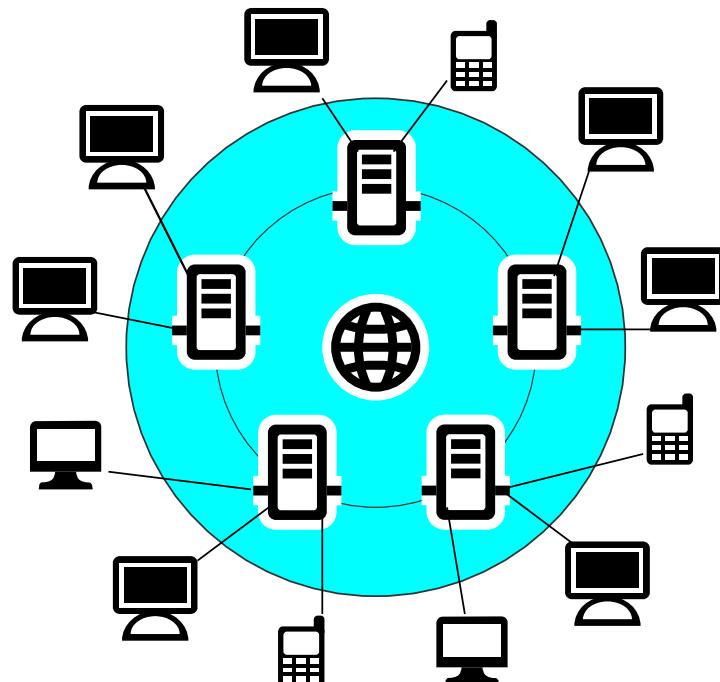
Desirable Experiment Properties



Live Network Experiments

Experimenting in a deployed distributed system

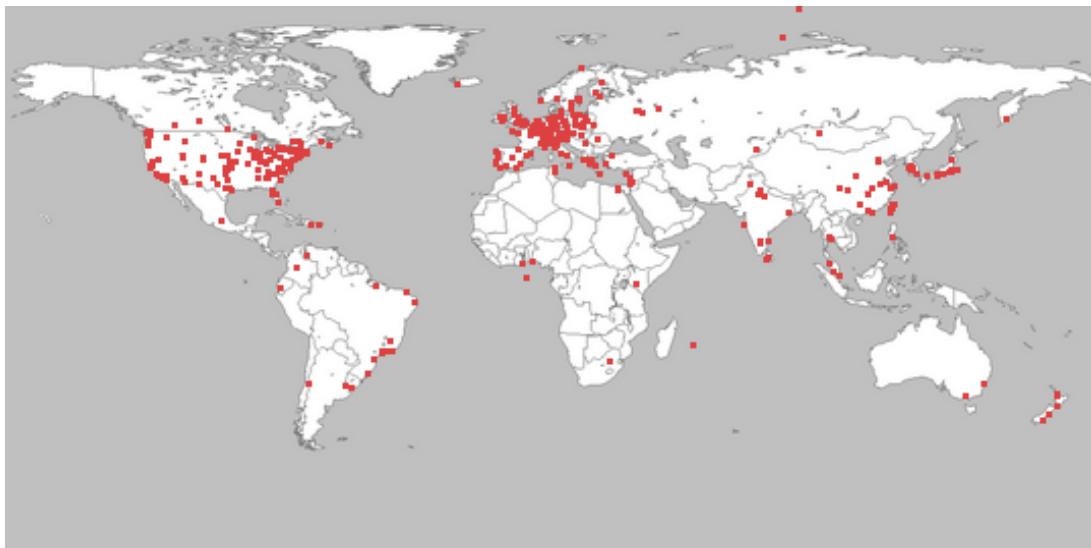
- Pros
 - Most realistic
 - The target environment
- Cons
 - Hard to manage/debug
 - Lengthy deployment
 - Security risks



Testbed Network Experiments

Experimenting in a distributed testbed (e.g. PlanetLab)

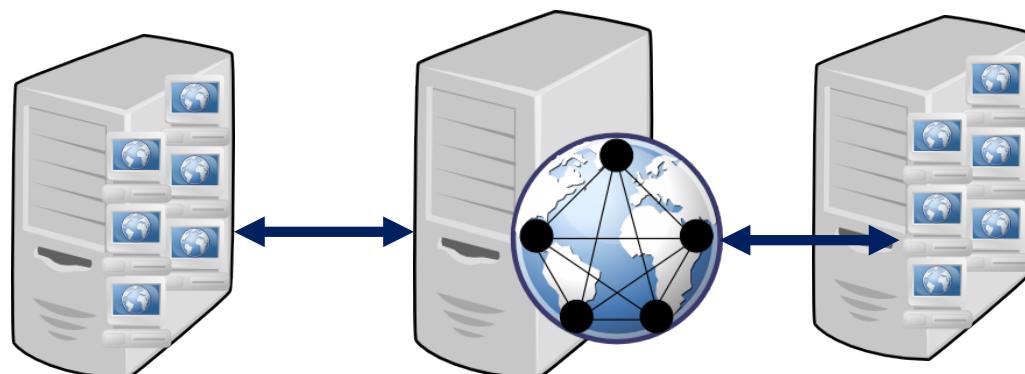
- Pros
 - Close to target environment
 - Runs on the Internet or uses Internet protocols
- Cons
 - Hard to manage and debug
 - Doesn't scale well in low-resource environs
 - Can be hard to model network properties



Network Emulation

Experimenting with network emulators (e.g. Modelnet)

- Pros
 - Runs on the target OS and uses Internet protocols
 - Can model various network properties
- Cons
 - Requires multiple machines and custom installed OS
 - Must run in real time
 - Per-process overhead may cause kernel issues



Network Simulation

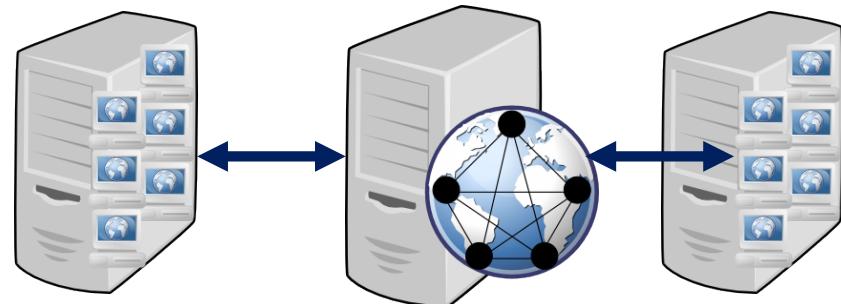
Experimenting with network simulators (e.g. NS3)

- Pros
 - Deterministic (reproducible)
 - Can model various network properties
 - Can scale very well
- Cons
 - Application model can be too abstract
 - Abstractions can lead to inaccurate results



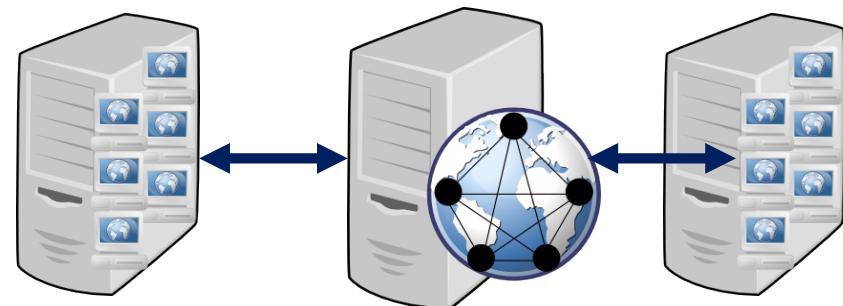
Simulation vs. Emulation: Realism

Simulation	Emulation
Abstracts away most system components	Runs the real OS, kernel, protocols, applications
Simulator is generally only internally consistent	Software is interoperable with external components
Less resource intensive	More resource intensive



Simulation vs. Emulation: Time

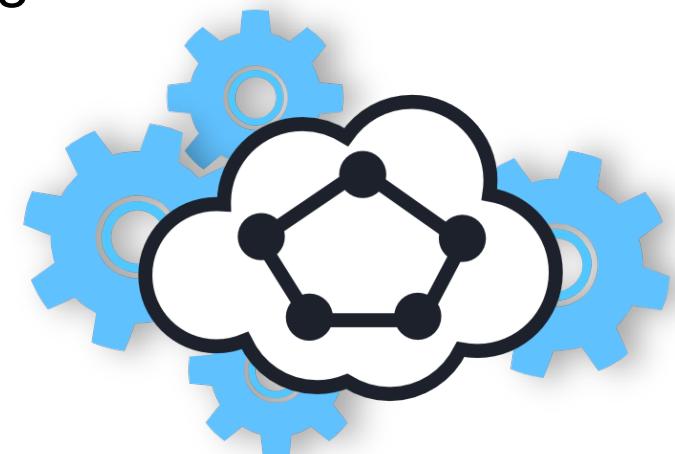
Simulation	Emulation
“As-fast-as-possible”	Real time
Control over clock, can pause time without issue	Time must advance in synchrony with wall-clock
Weak hardware extends total experiment runtime	Weak hardware causes glitches that are difficult to detect and diagnose



Shadow Design

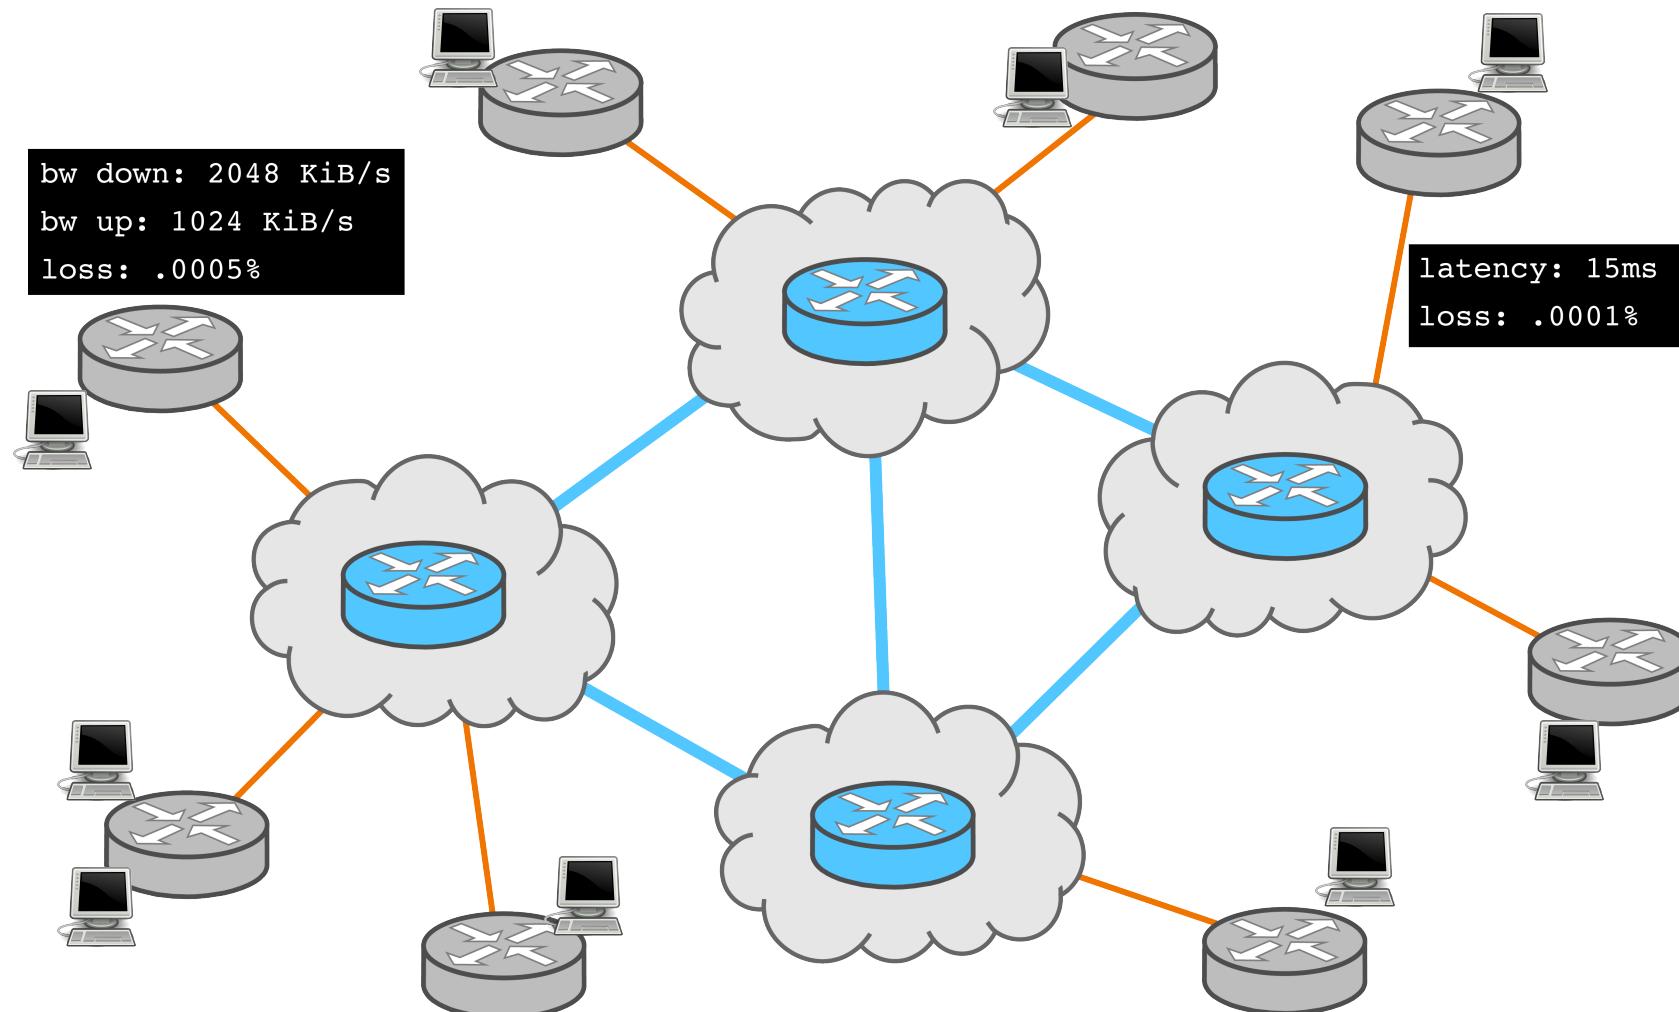
Shadow Overview

- Parallel discrete-event network simulator
- Models routing, latency, bandwidth
- Simulates time, CPU, OS
 - TCP/UDP, sockets, queuing, threading
- Emulates POSIX C API on Linux
- Directly executes apps as plug-ins

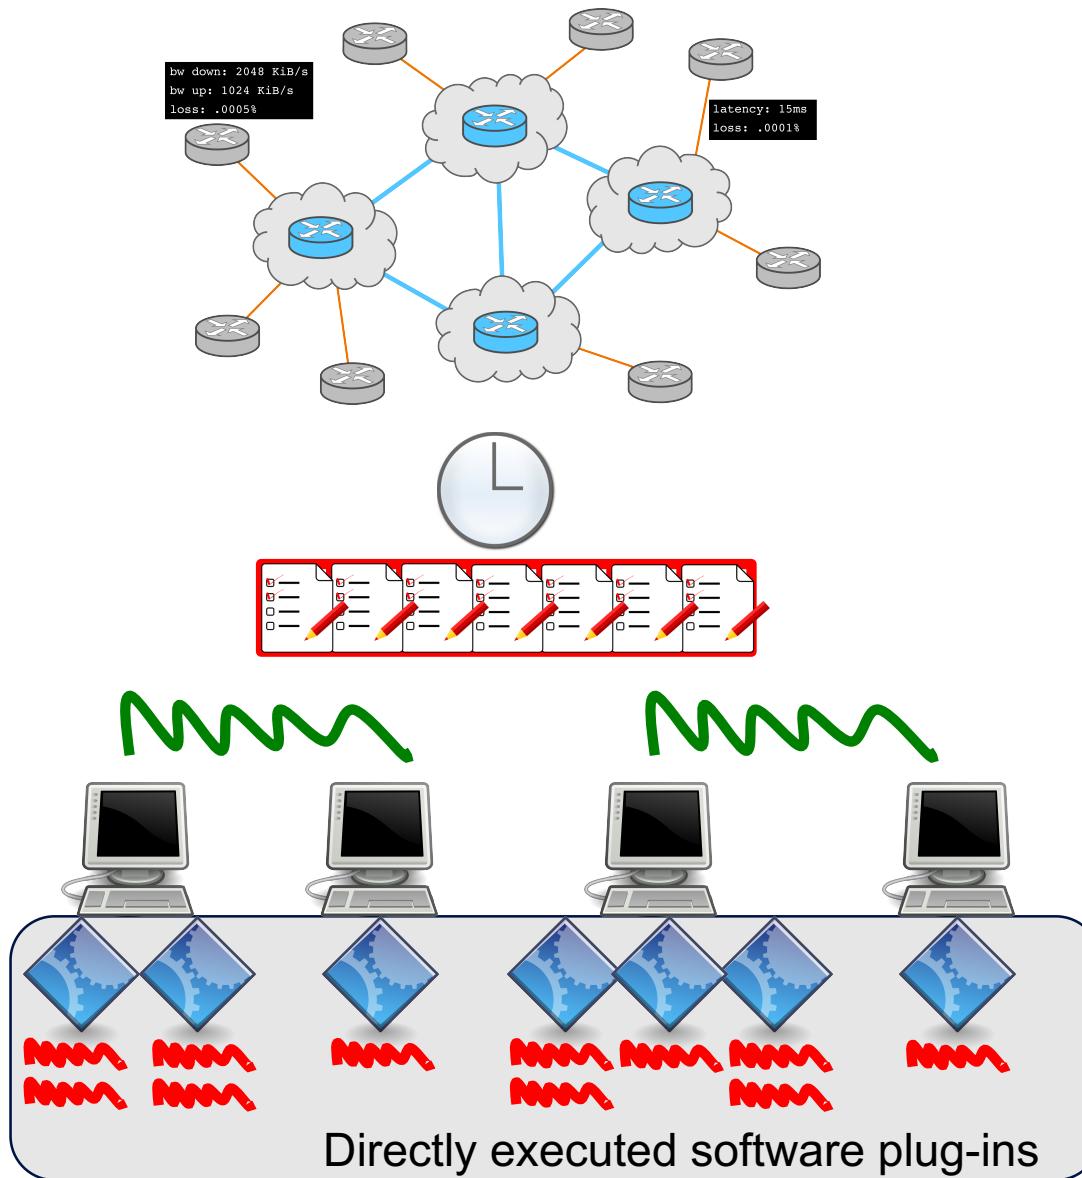


Initializing the Simulation

Load network model, create virtual hosts



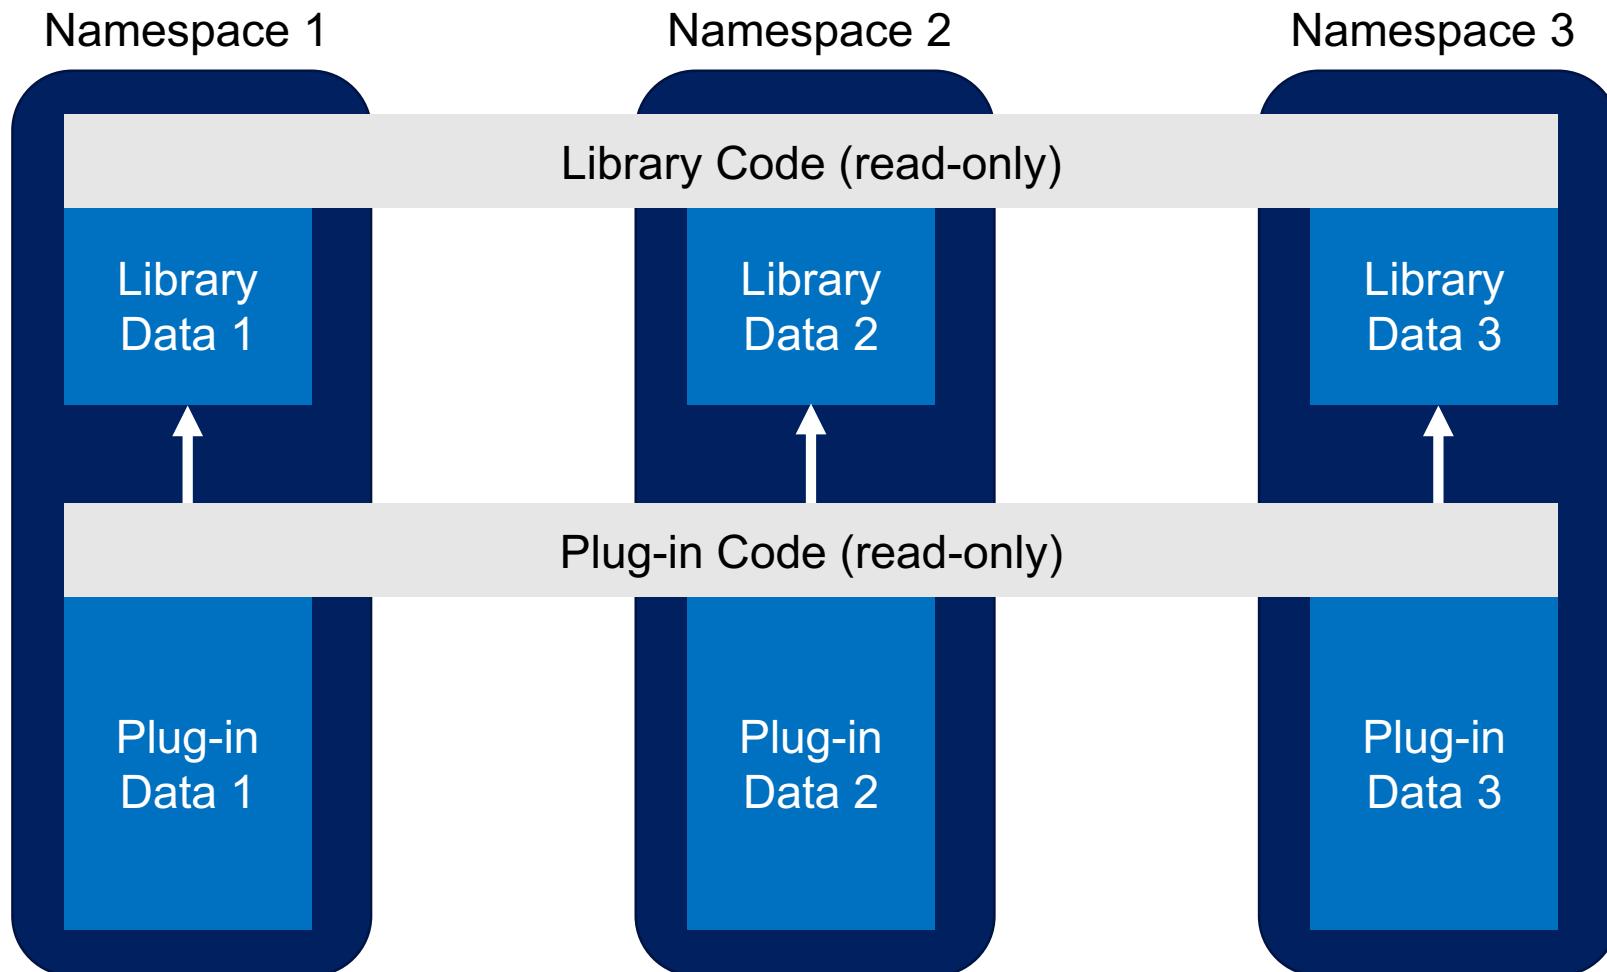
Shadow Simulation Layout



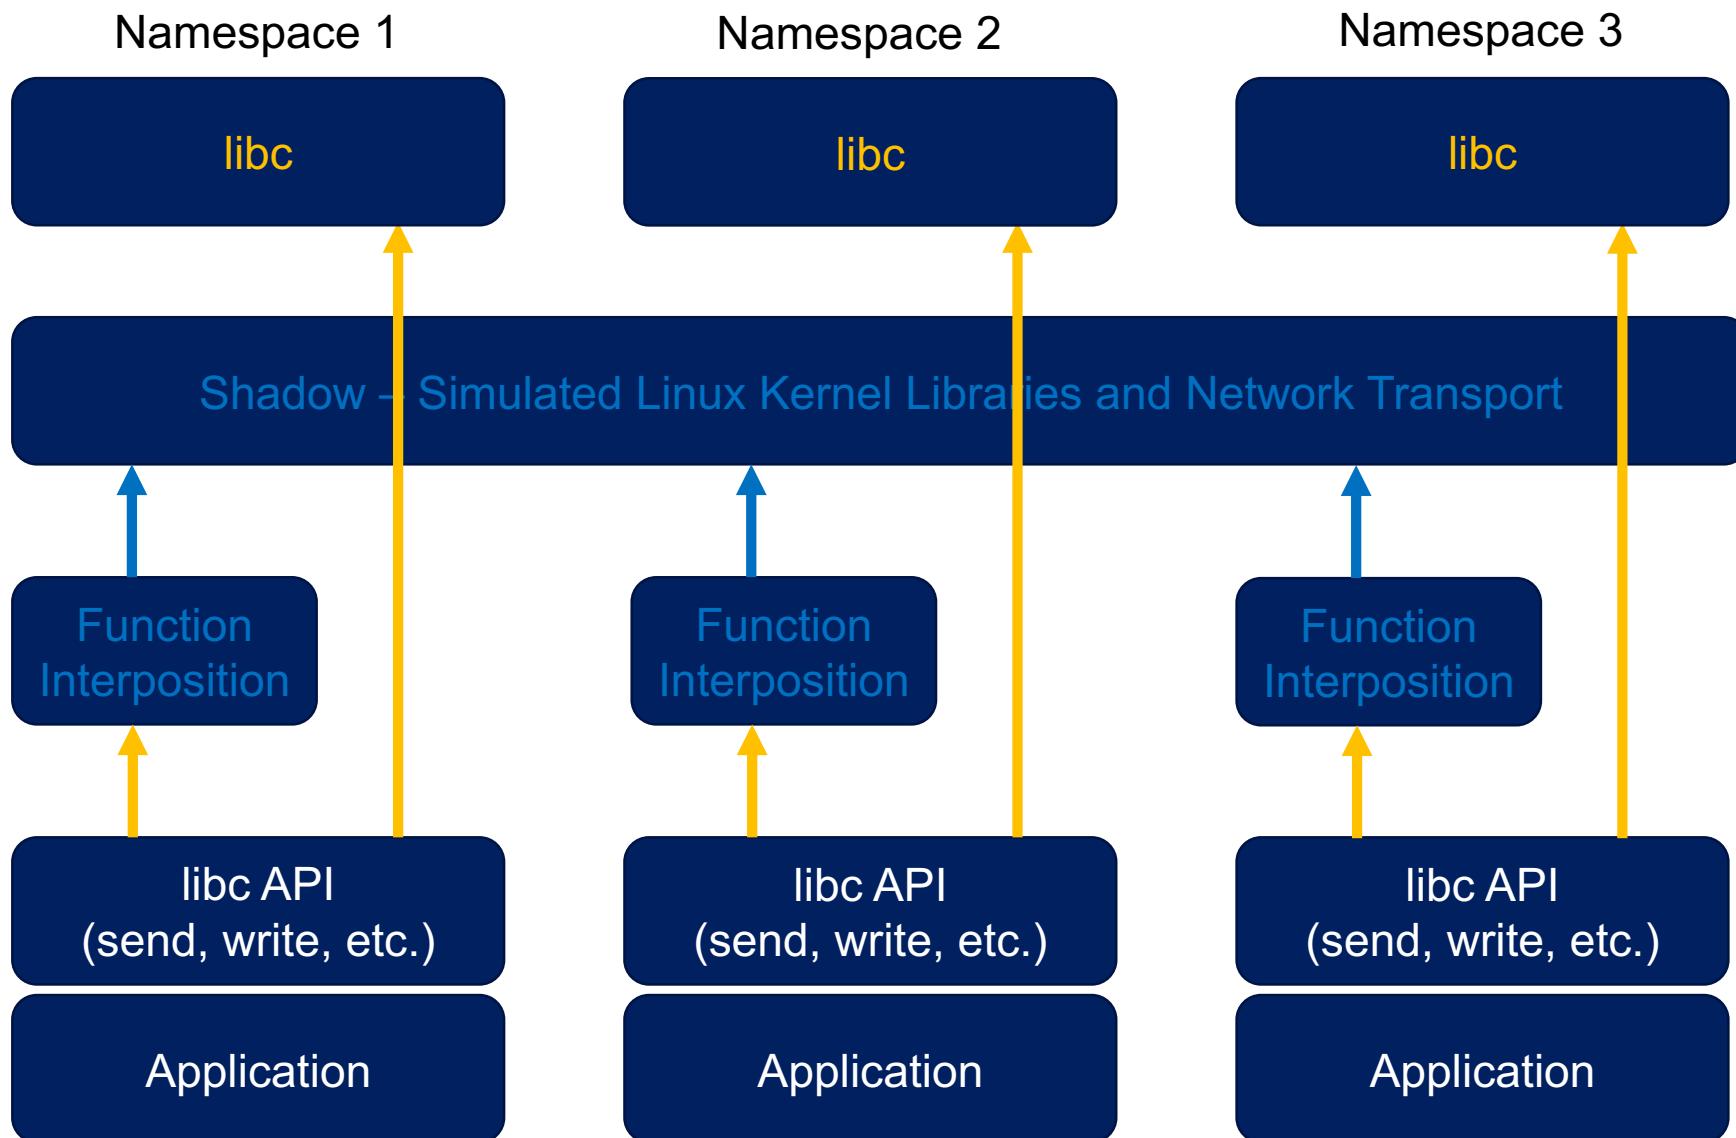
- Network model
- Global simulation clock
- Discrete event queue
- Shadow worker threads
- Virtual hosts
- Virtual processes (i.e. namespaces)
- Virtual threads

App Memory Management

Apps loaded in independent namespaces, copy-on-write



Direct Execution in a Simulator

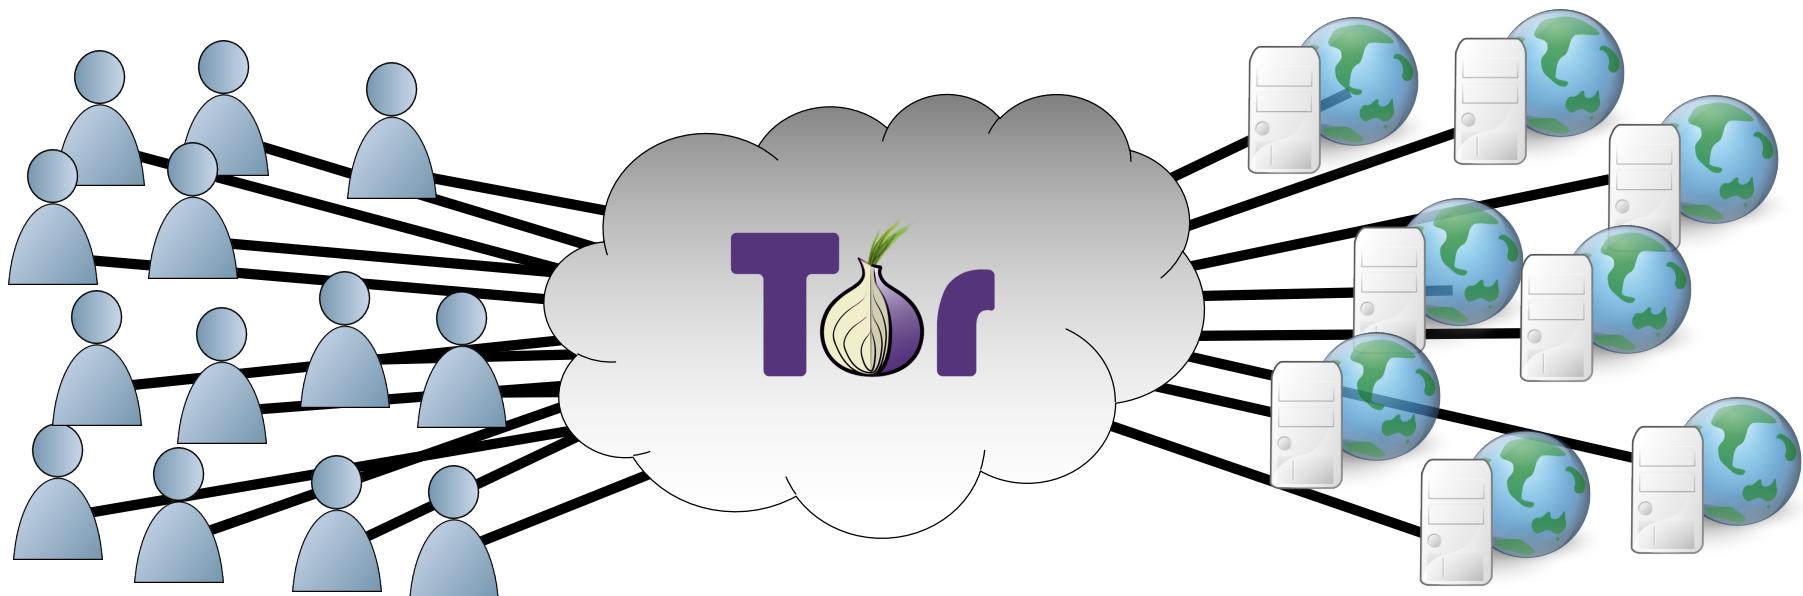


Simulation Use Cases – The Tor Anonymity Network

Gentle Tor Introduction

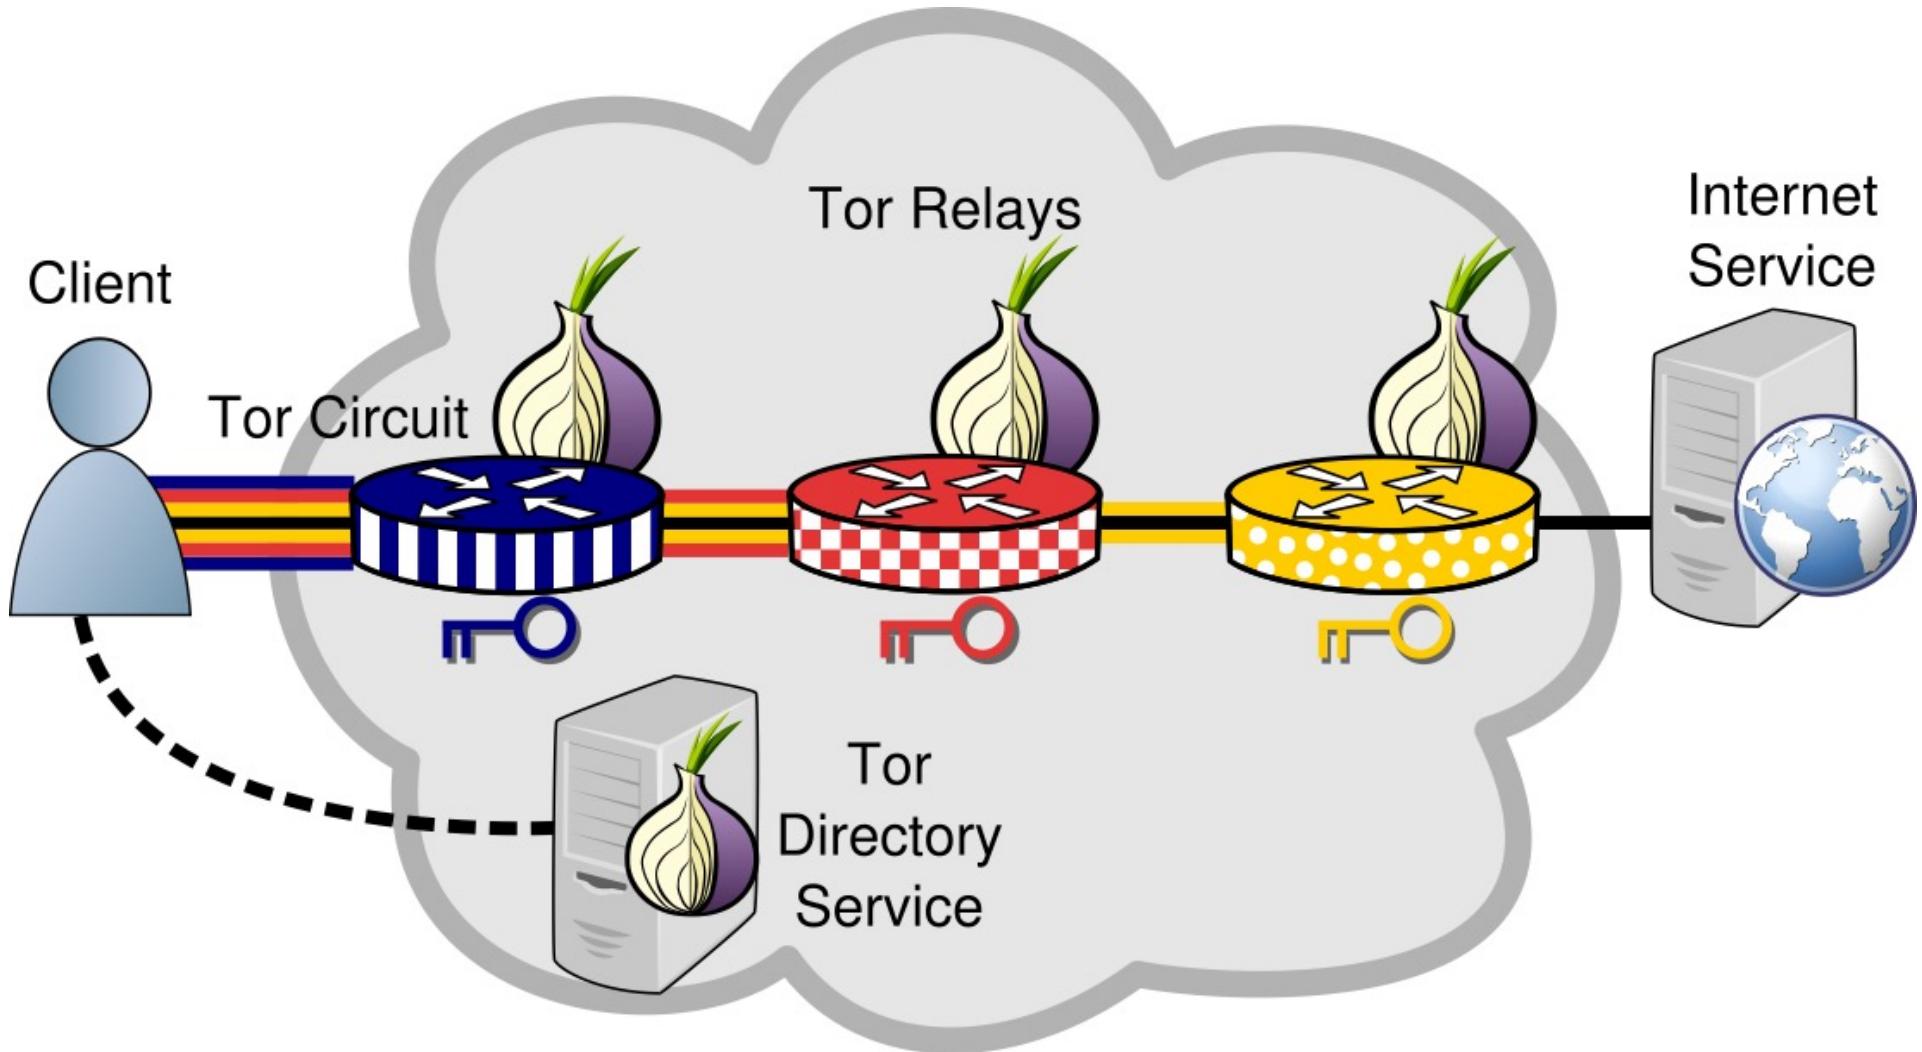
Tor Overview

Tor: a censorship resistant, privacy-enhancing anonymous communication system



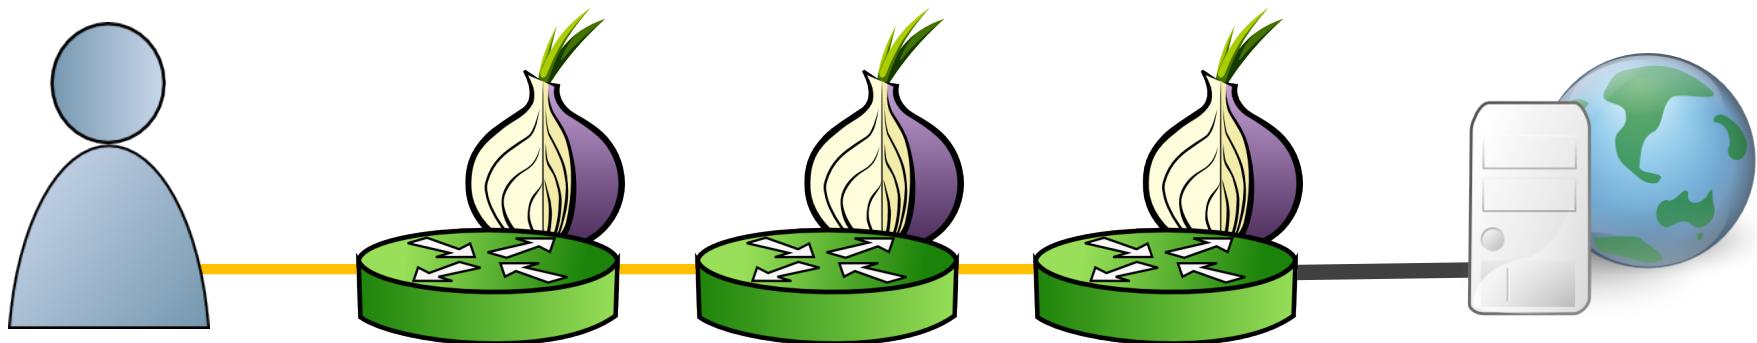
Estimated ~1.75 M. Users/Day
(metrics.torproject.org)

Anonymous Communication

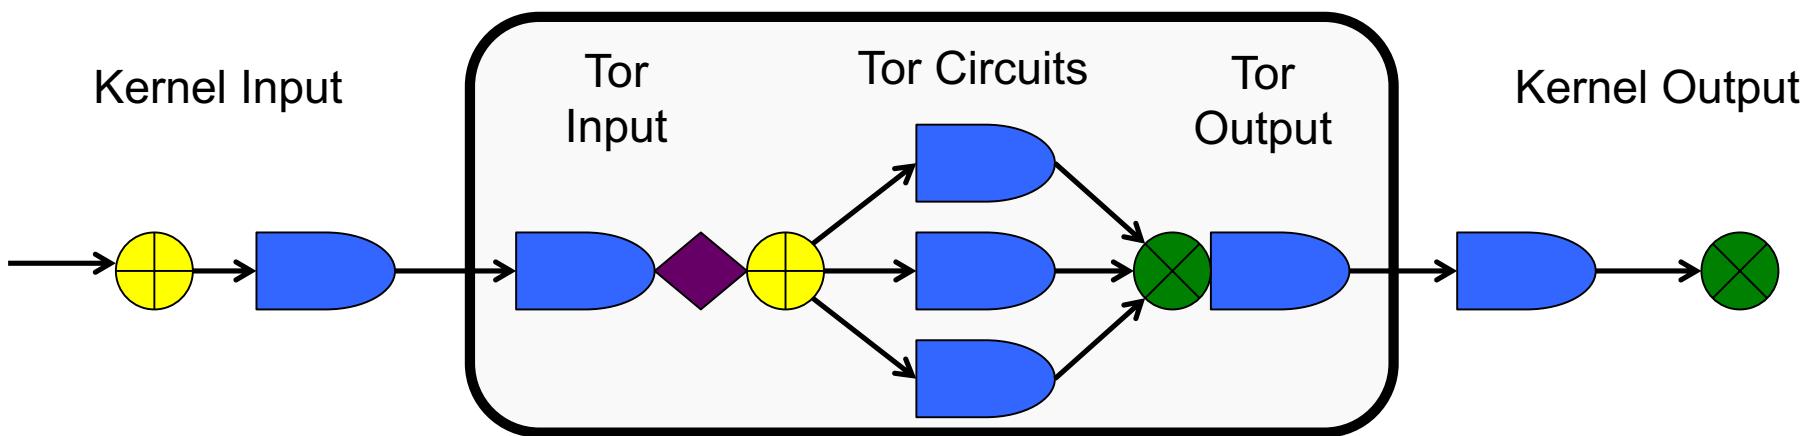
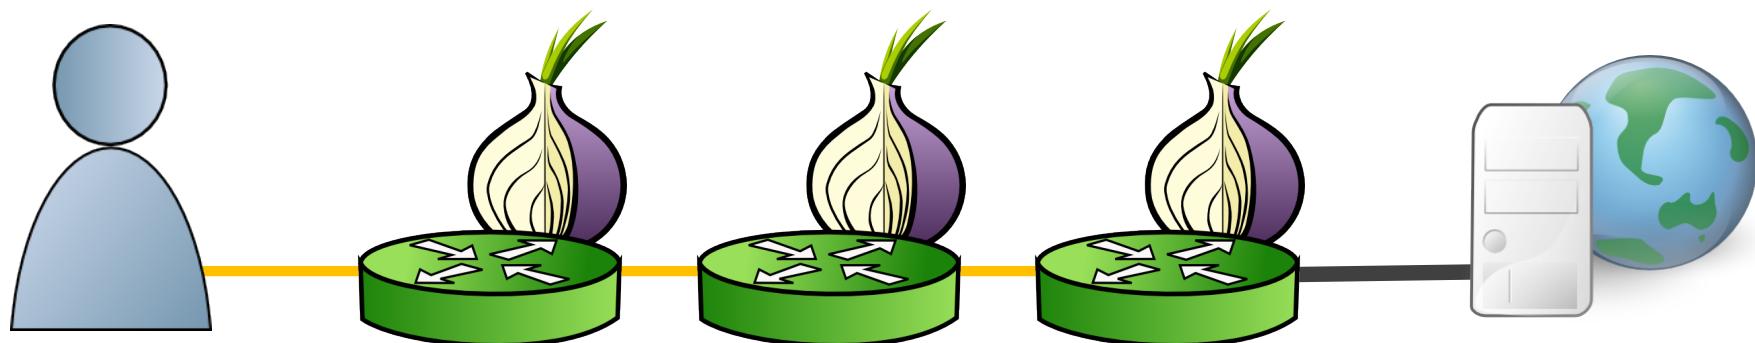


Congestion Analysis

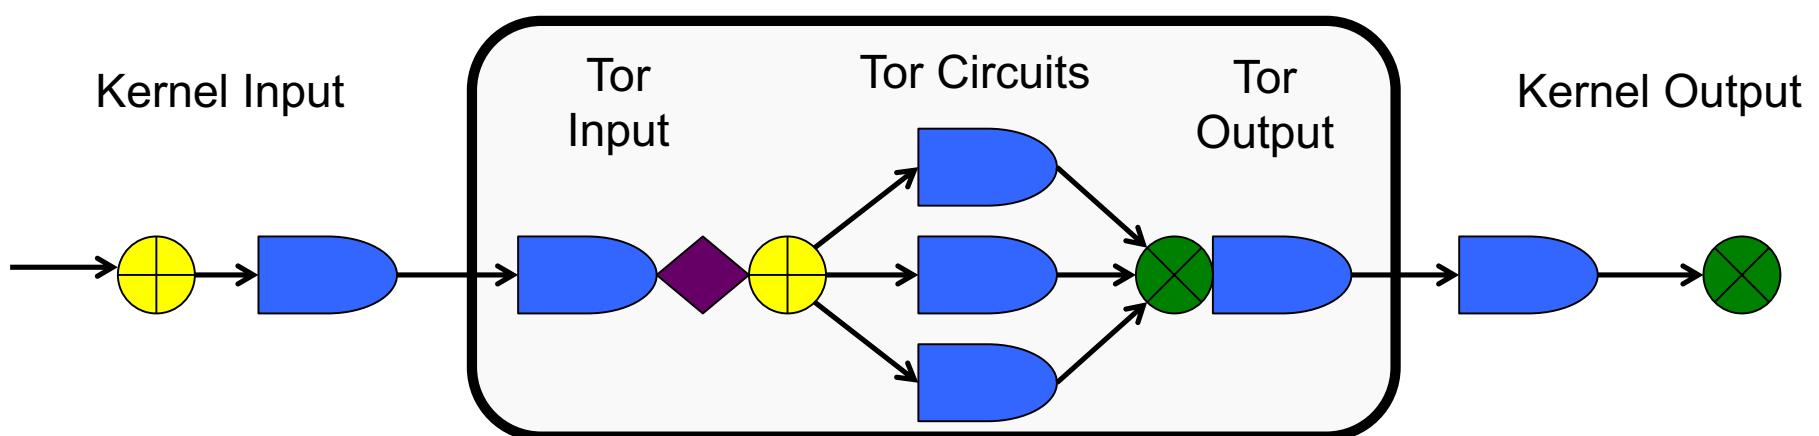
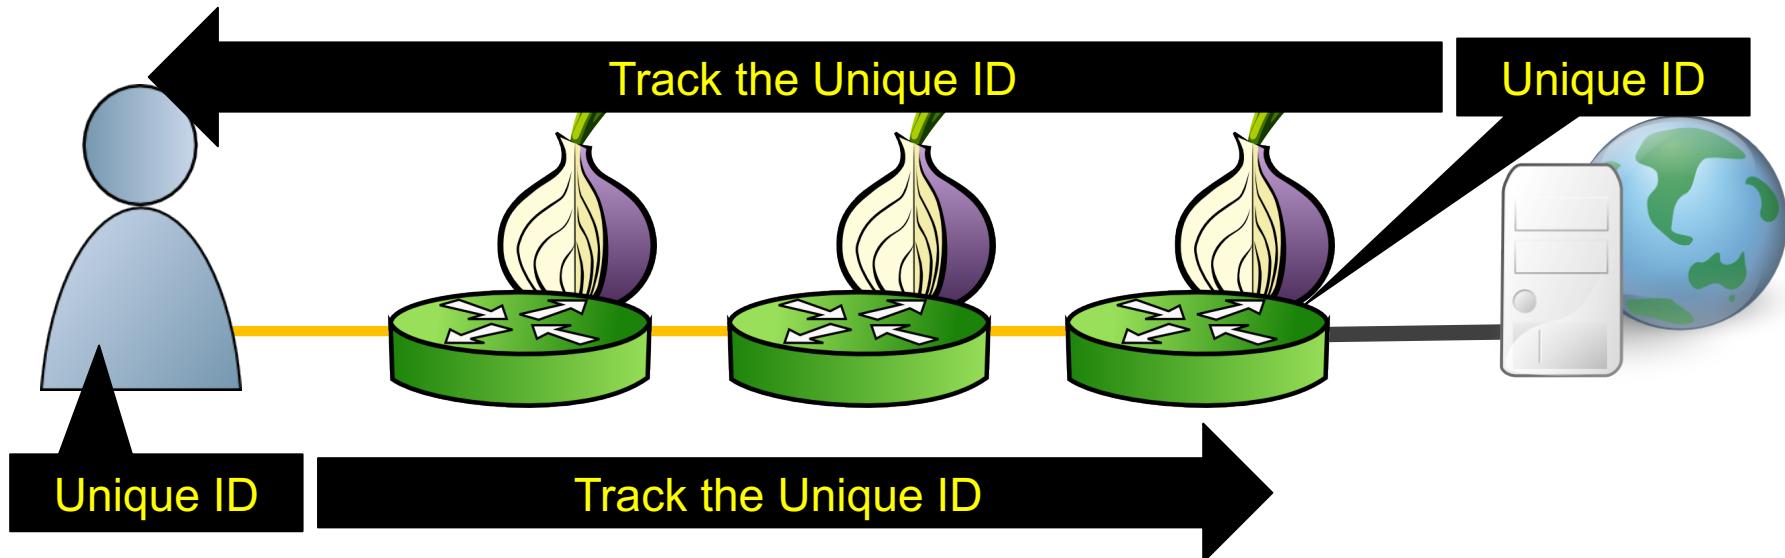
Multiple Hops – Tor is Slower



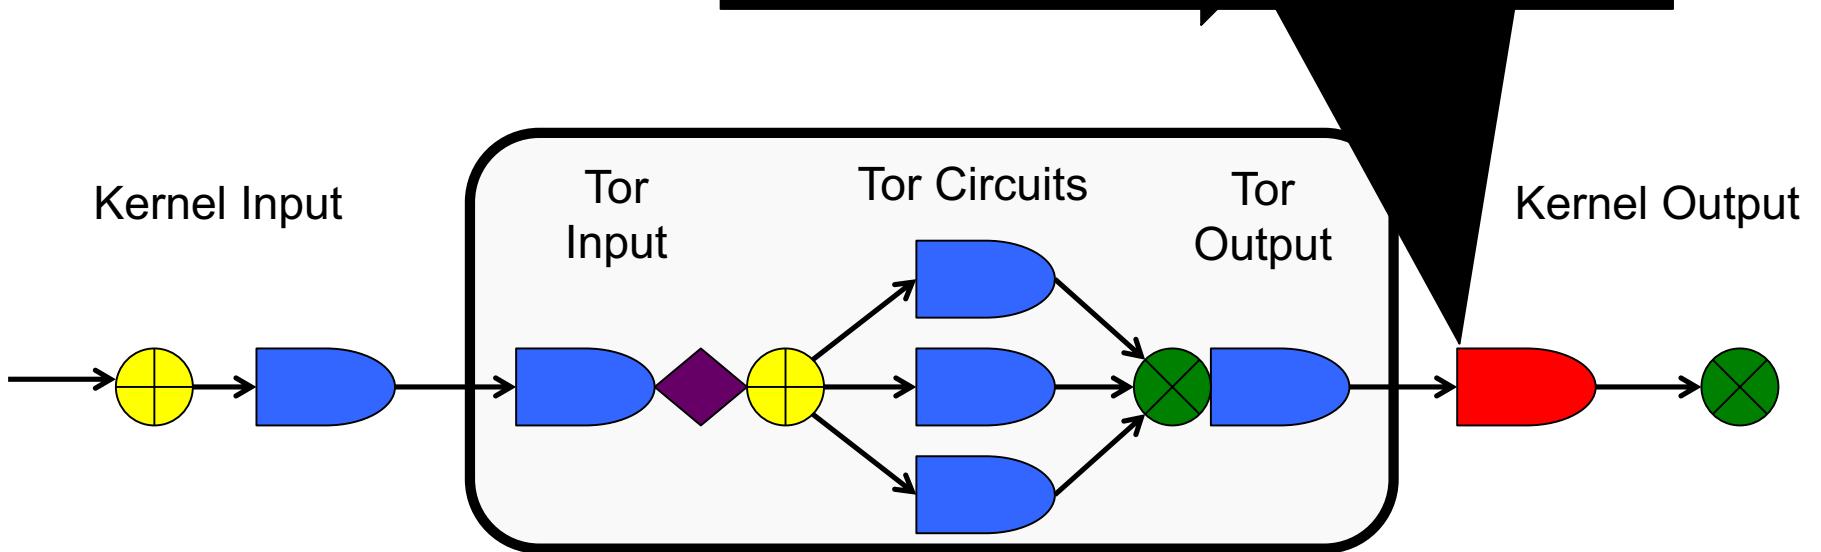
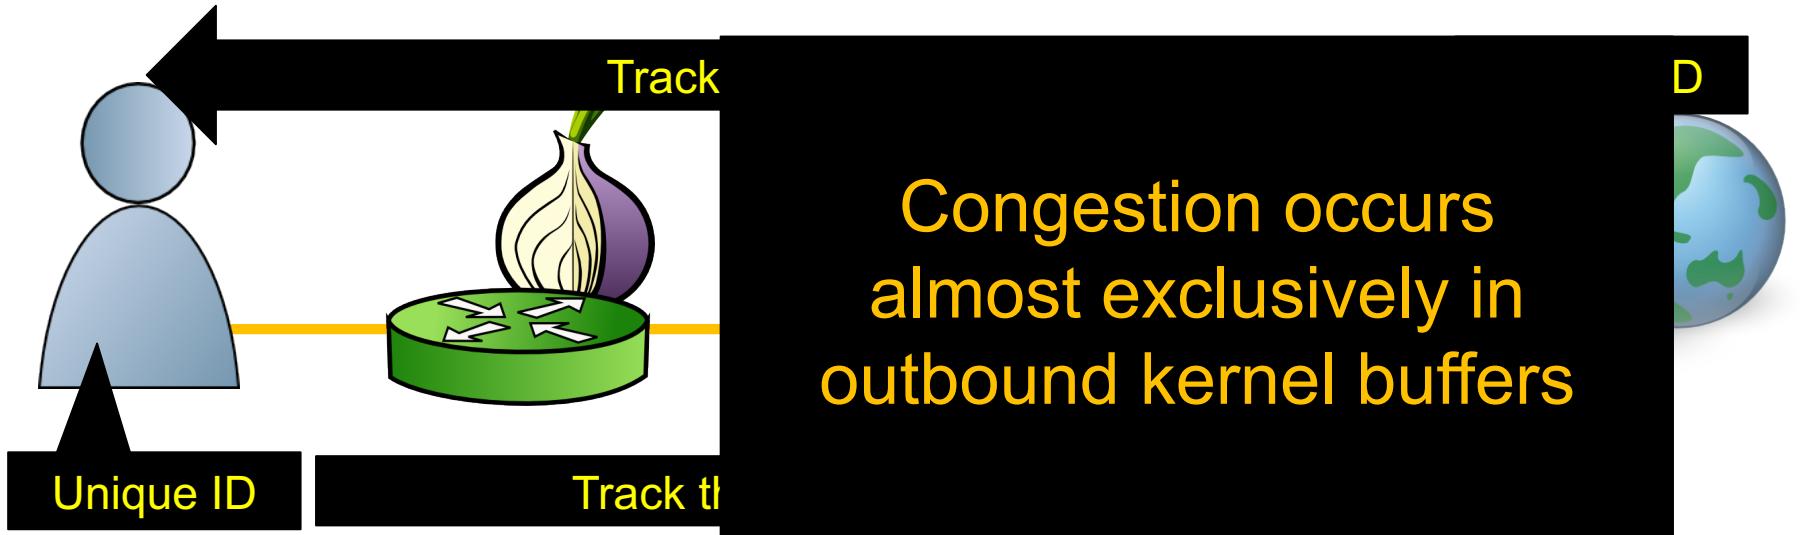
Buffers in a Tor Relay



Tracking Congestion in Tor



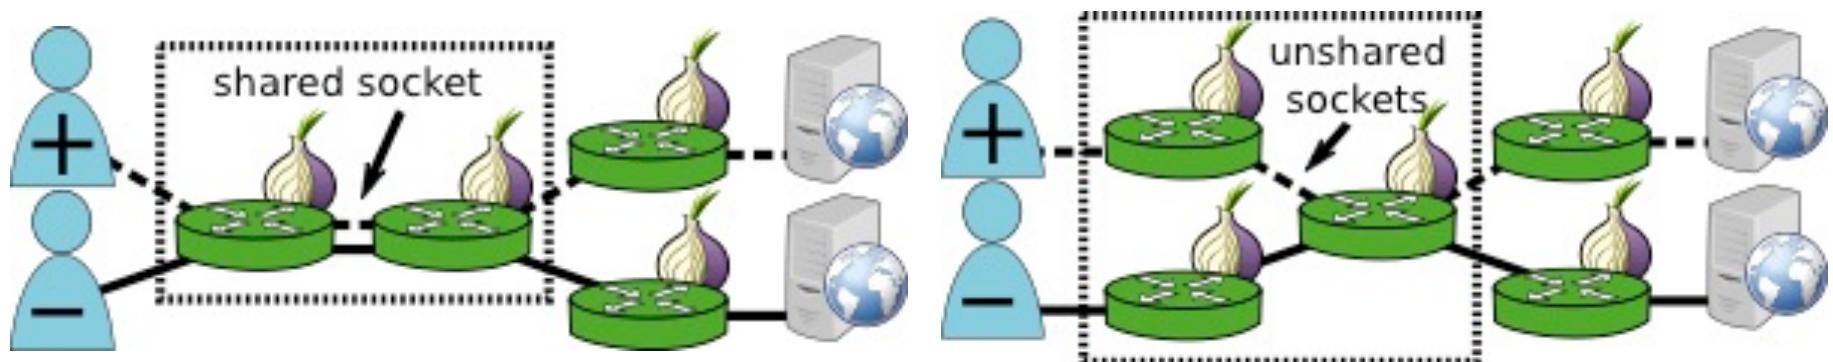
Tracking Congestion in Tor



Scheduling Analysis

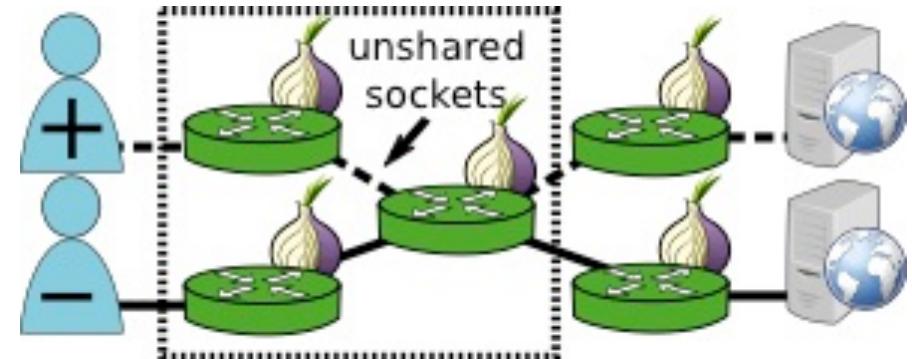
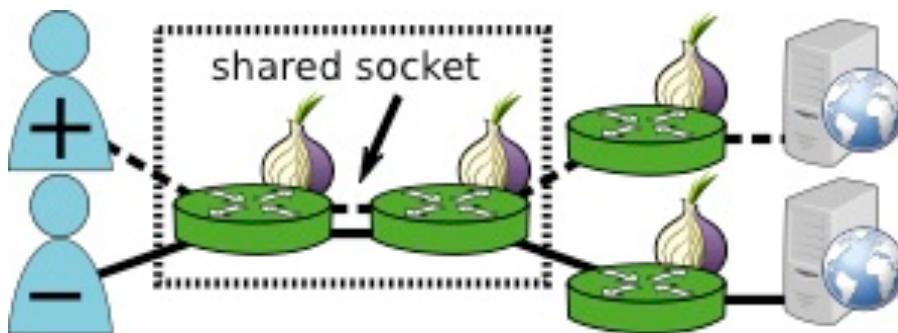
Tor Circuit Priority

Scenarios to understand outbound queue congestion

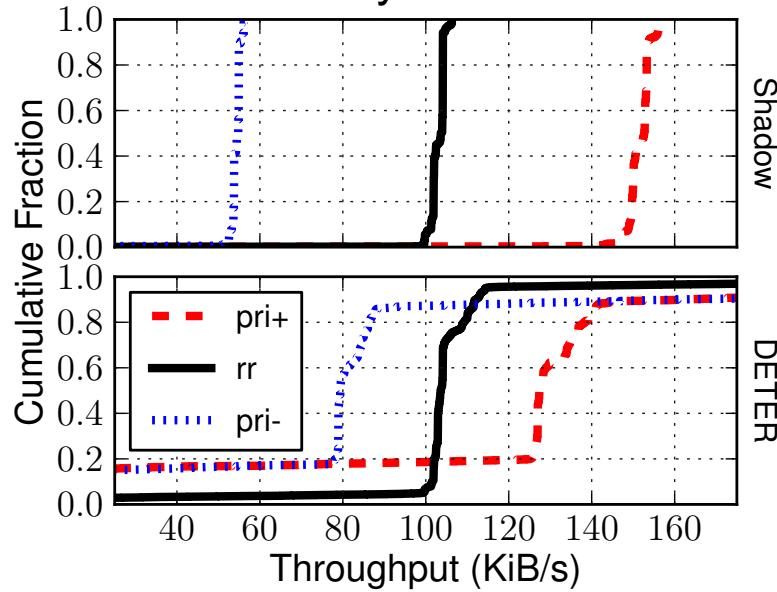


Tor Circuit Priority

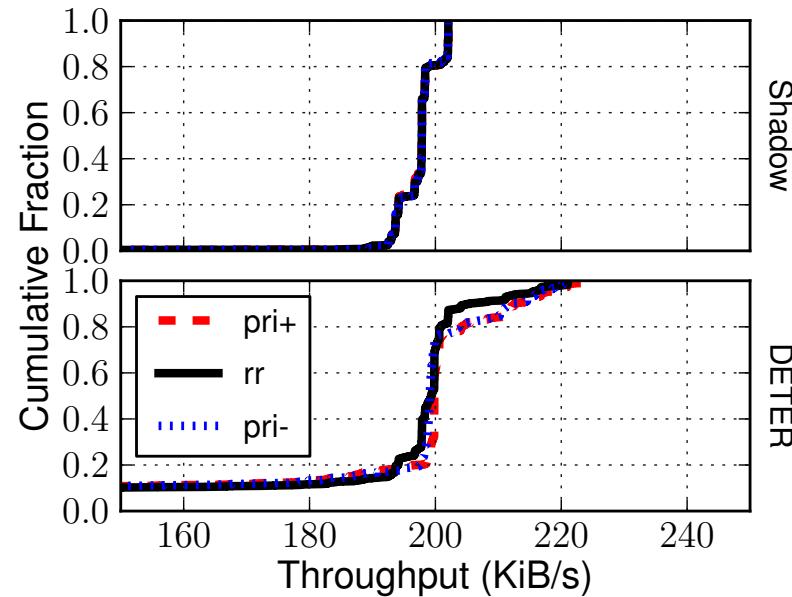
Scenarios to understand outbound queue congestion



Correctly differentiated

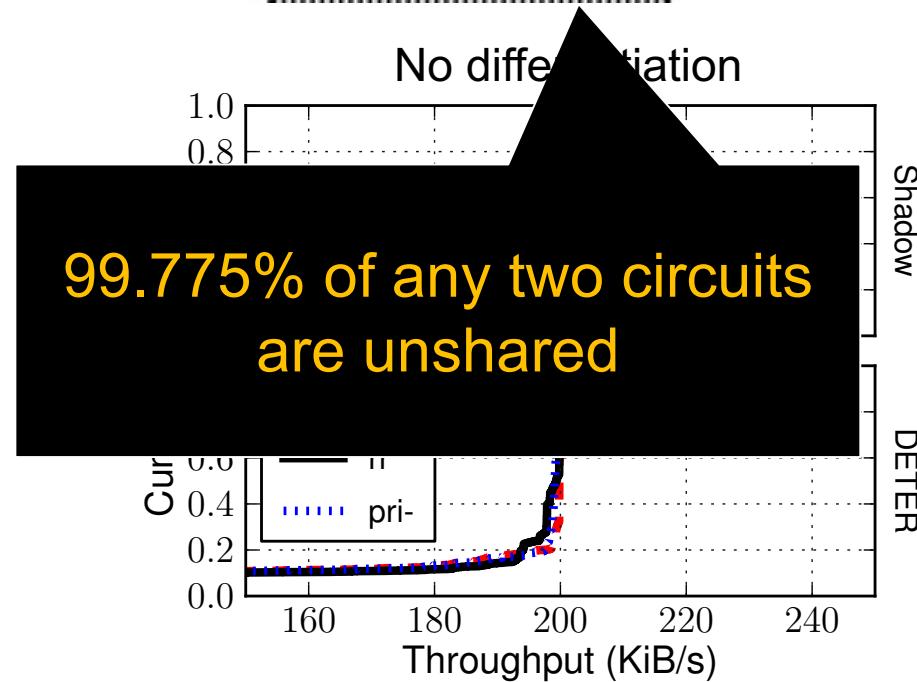
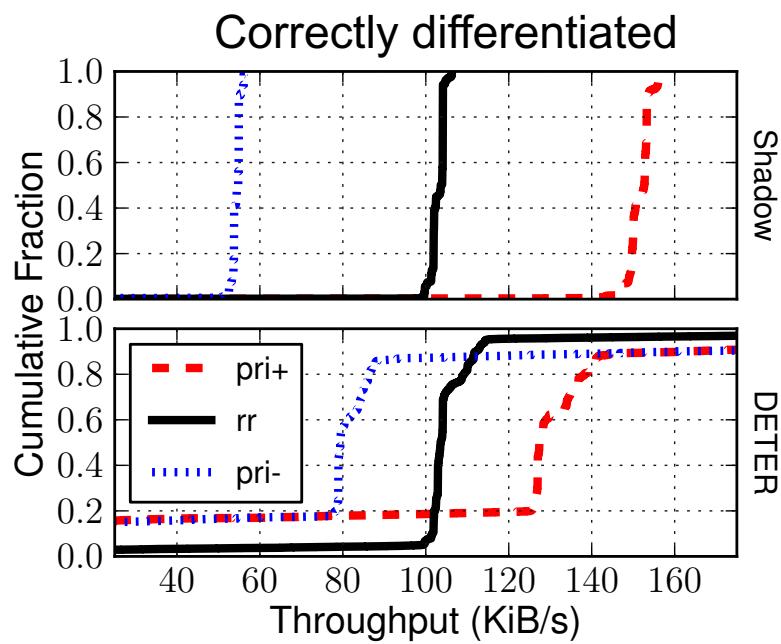
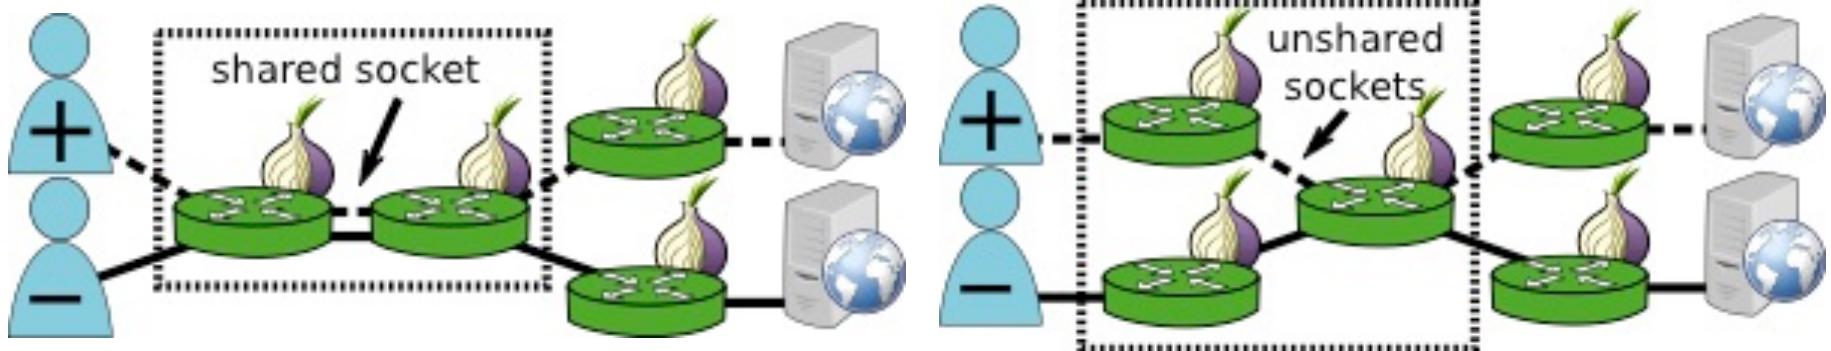


No differentiation



Tor Circuit Priority

Scenarios to understand outbound queue congestion



Performance Analysis

Queuing delays in kernel output buffer

Problem 1: Circuit scheduling

Solution:

- Don't handle sockets individually
- Process flush requests from all writable sockets

Problem 2: Flushing to sockets – buffer bloat

Solution:

- Don't write if kernel can't send
- Use TCP info to bound kernel writes

Network Simulation with Shadow

Enhanced Shadow with several missing TCP algorithms

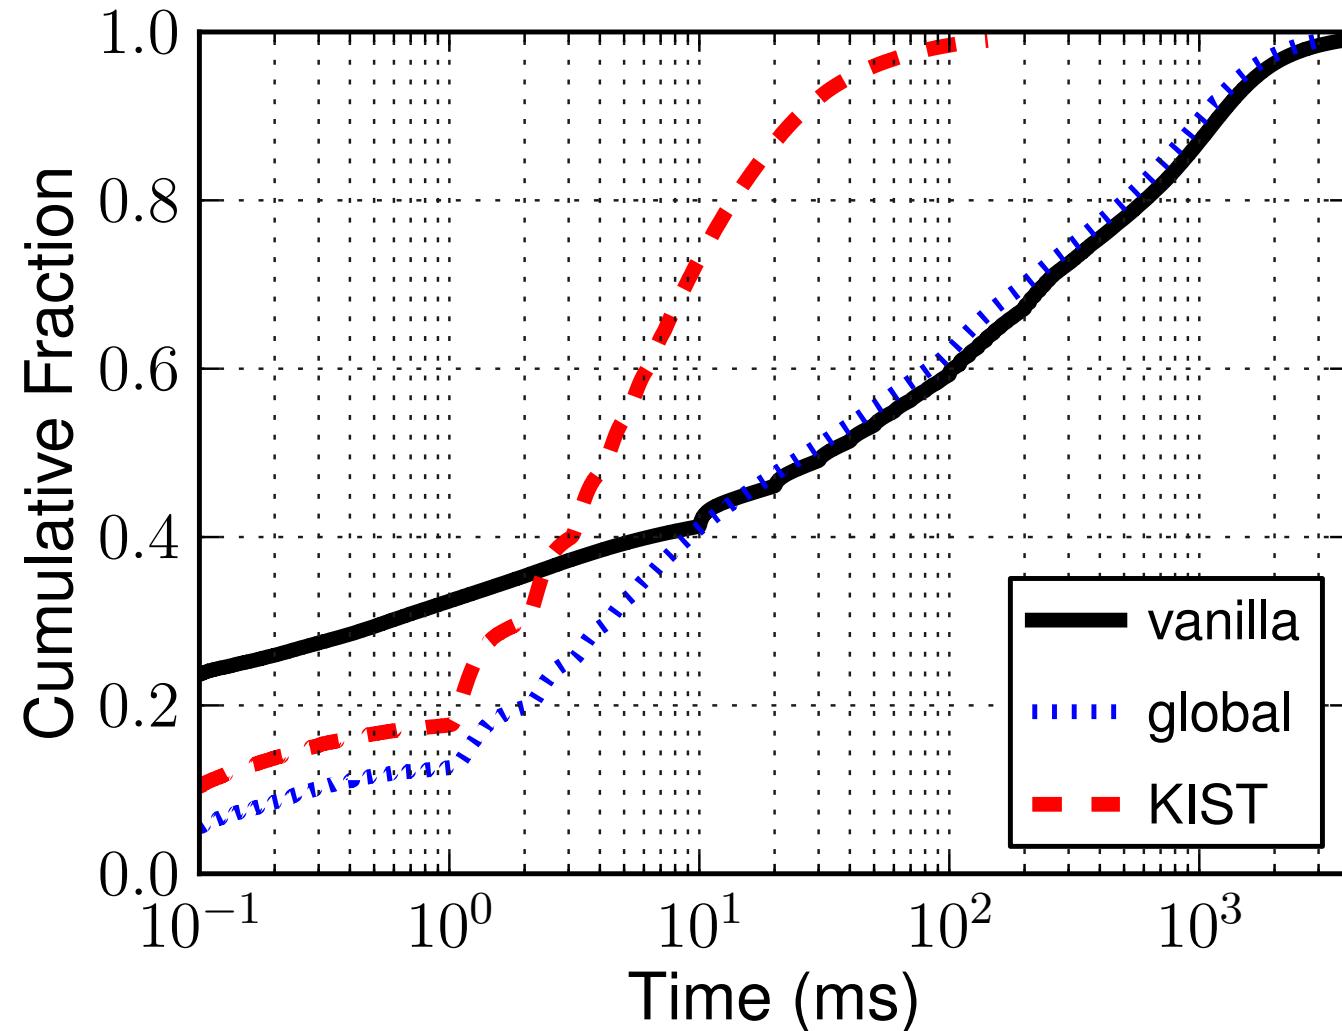
- CUBIC congestion control
- Retransmission timers
- Selective acknowledgements (SACK)
- Forward acknowledgements (FACK)
- Fast retransmit/recovery

Designed largest known private Tor network

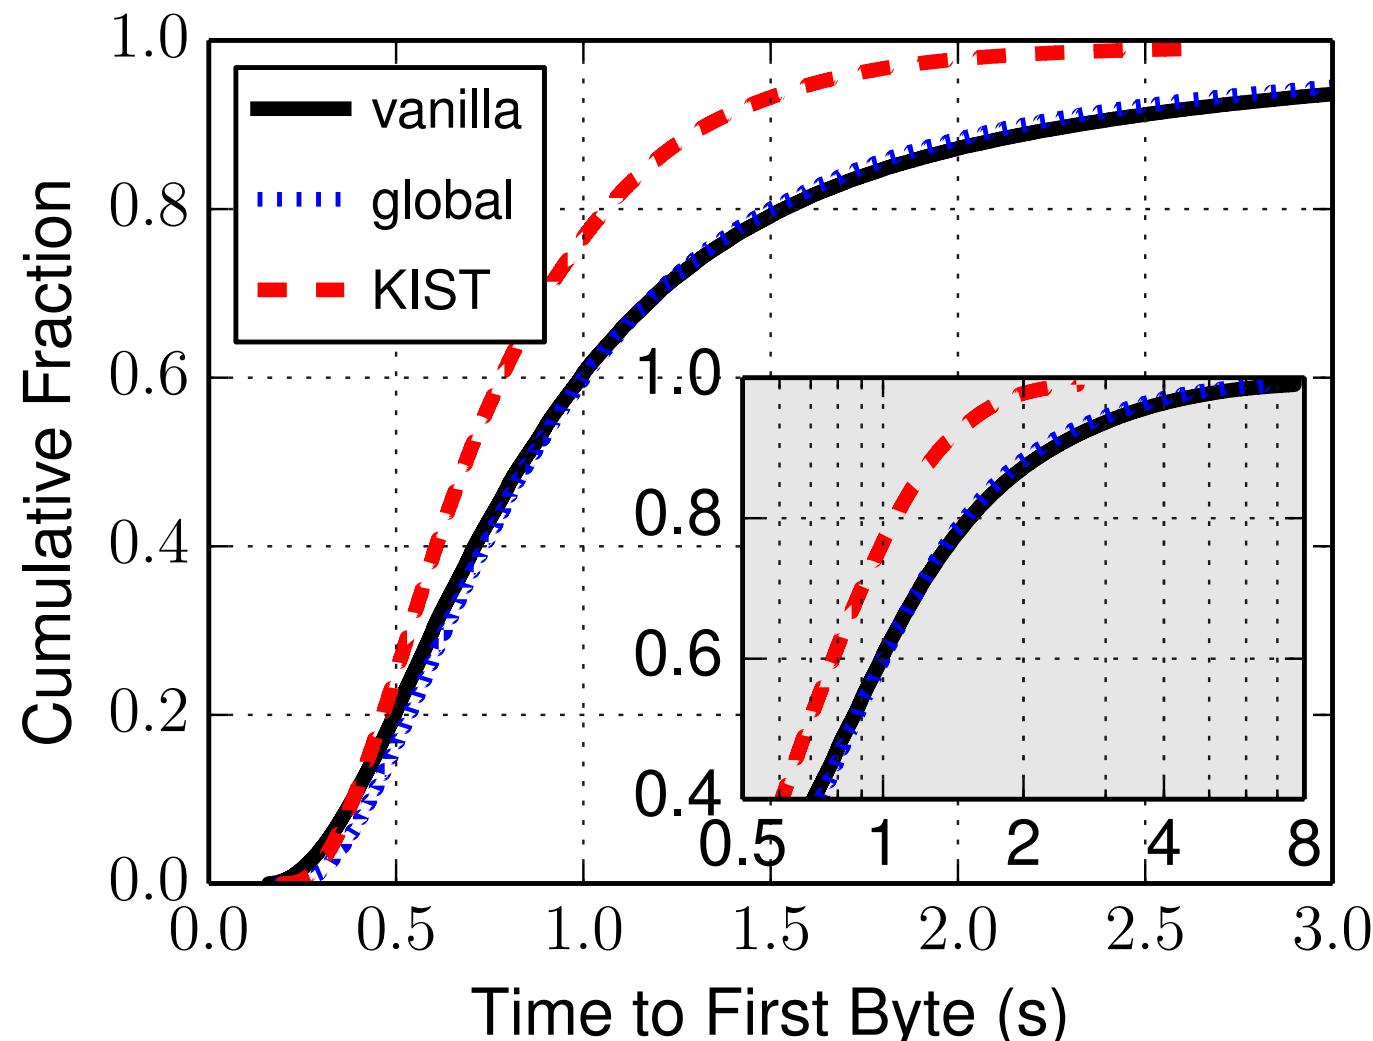
- 3600 relays and 12000 simultaneously active clients
- Internet topology graph: ~700k vertices and 1.3m edges

Analyze network-wide effects

KIST Reduces Kernel Congestion



KIST Improves Network Latency

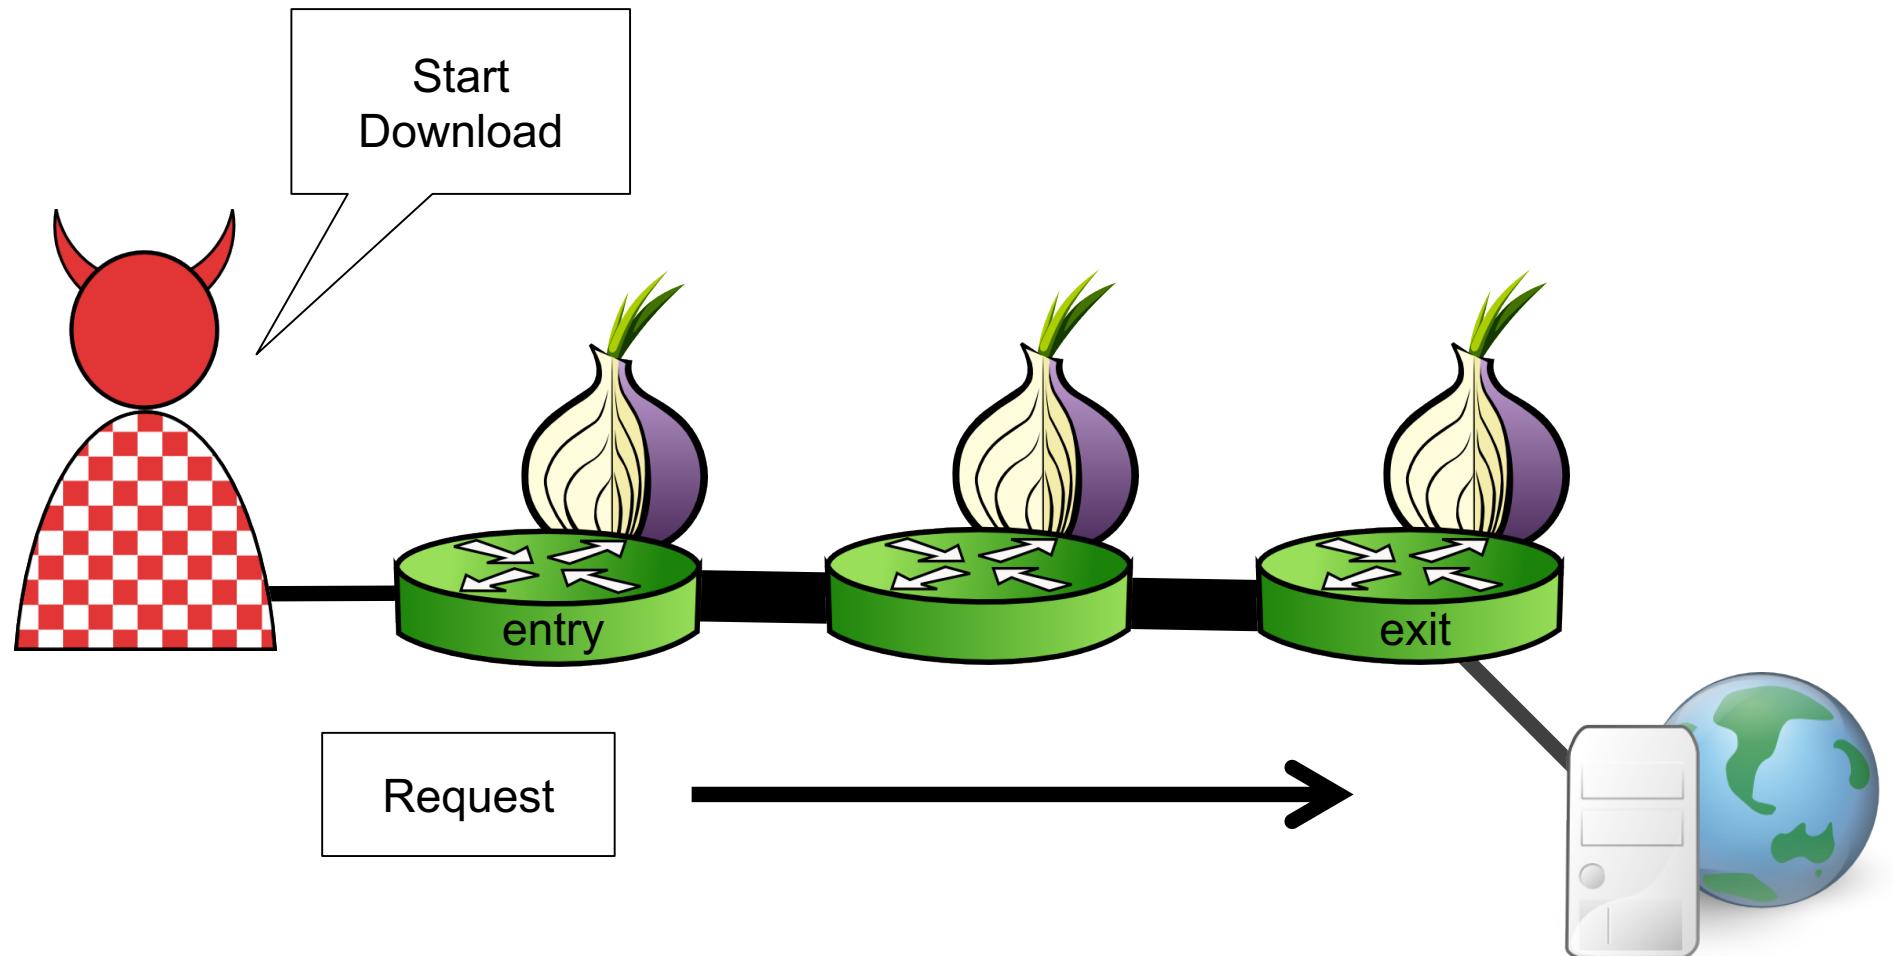


Denial of Service Attacks

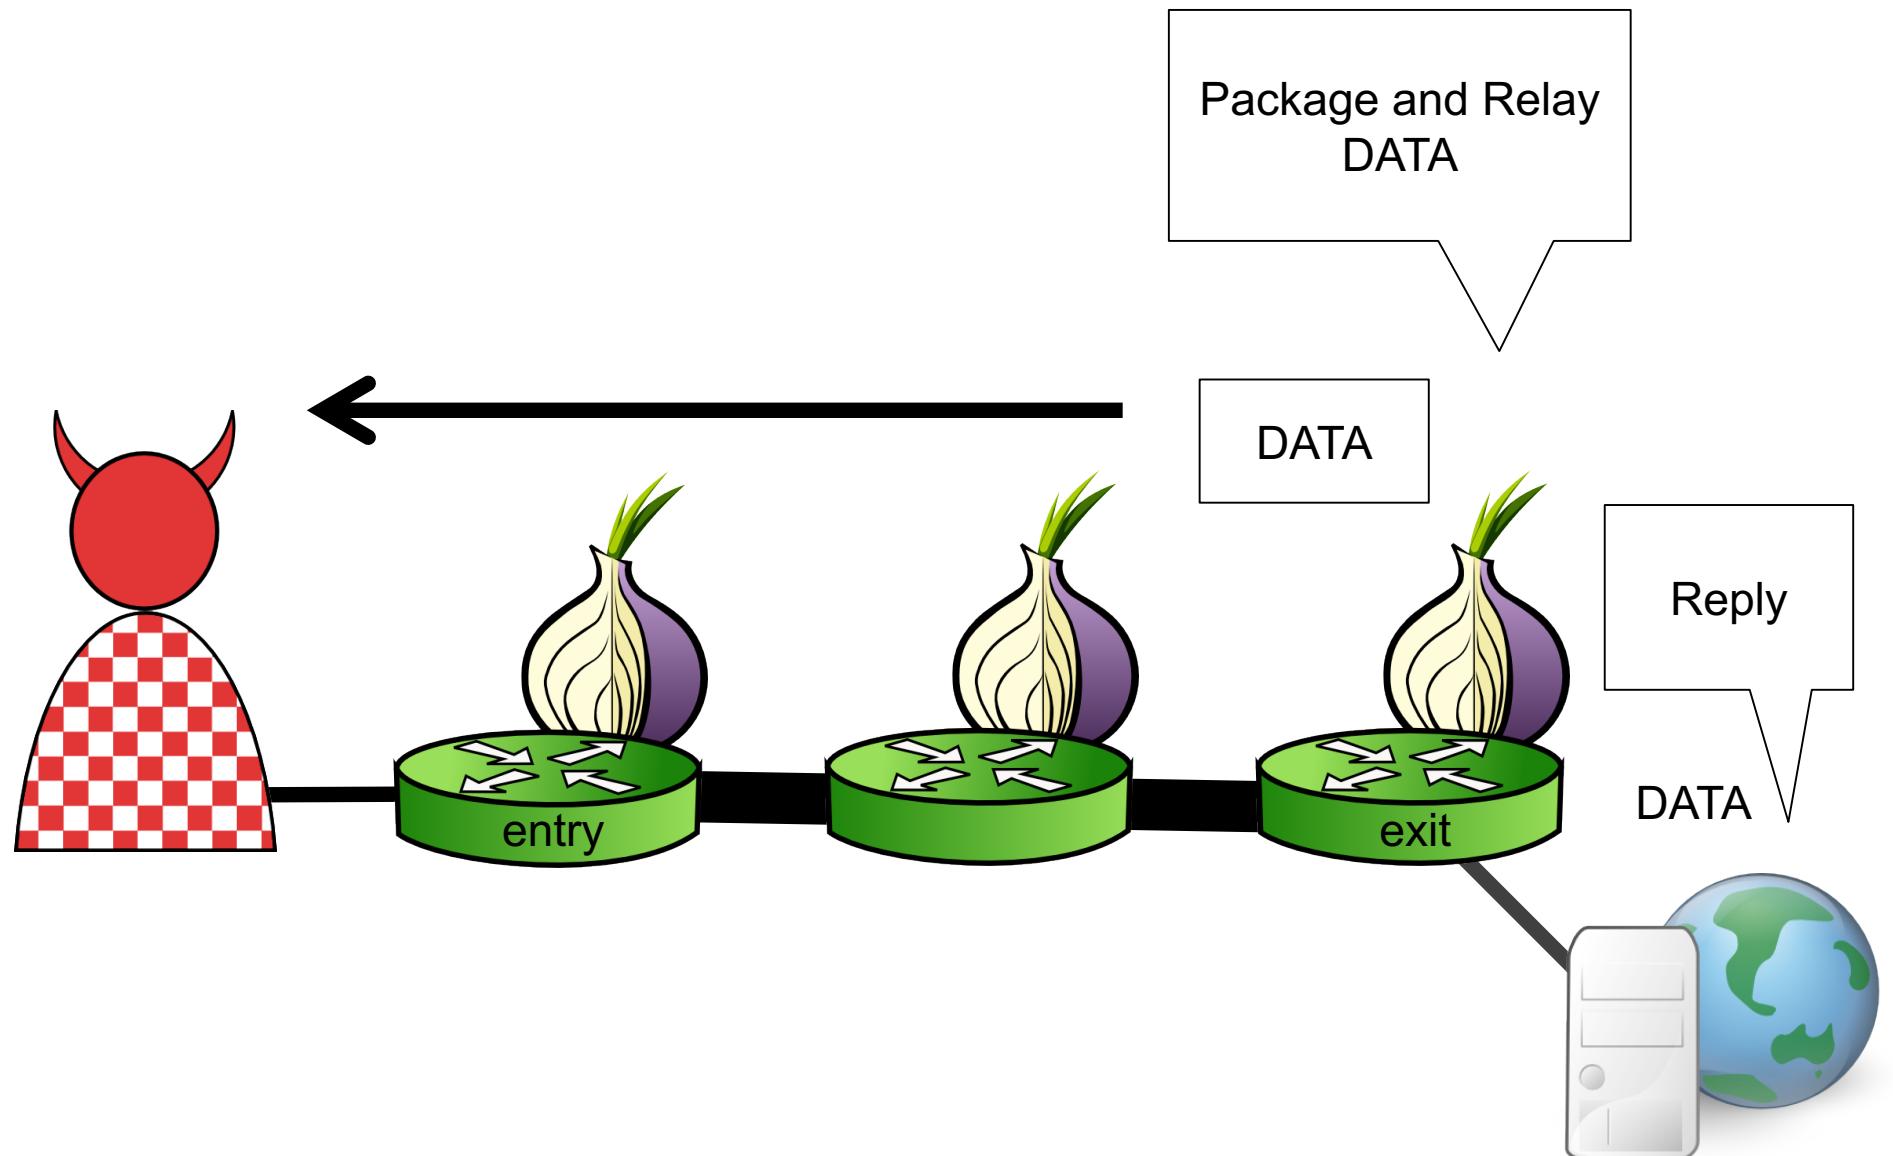
The Sniper Attack

- Memory-based denial of service (DoS) attack
- Exploits vulnerabilities in Tor's flow control protocol
- Can be used to disable arbitrary Tor relays

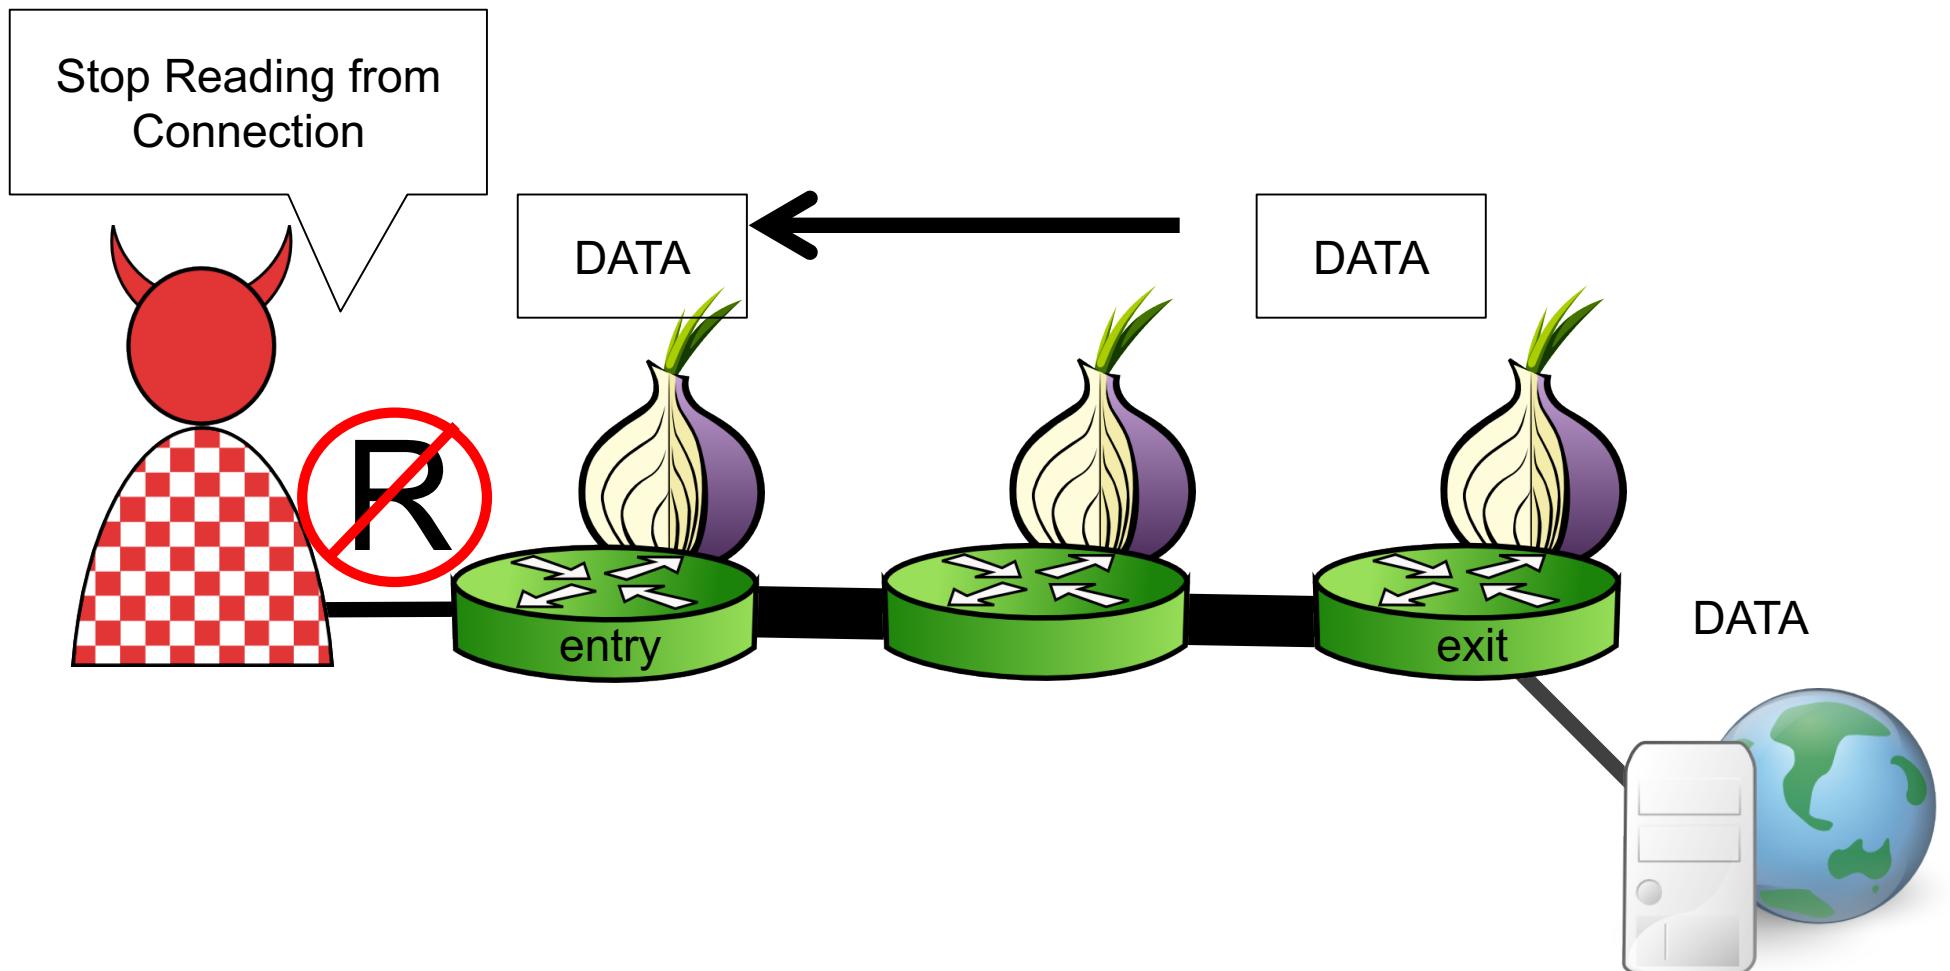
The Sniper Attack



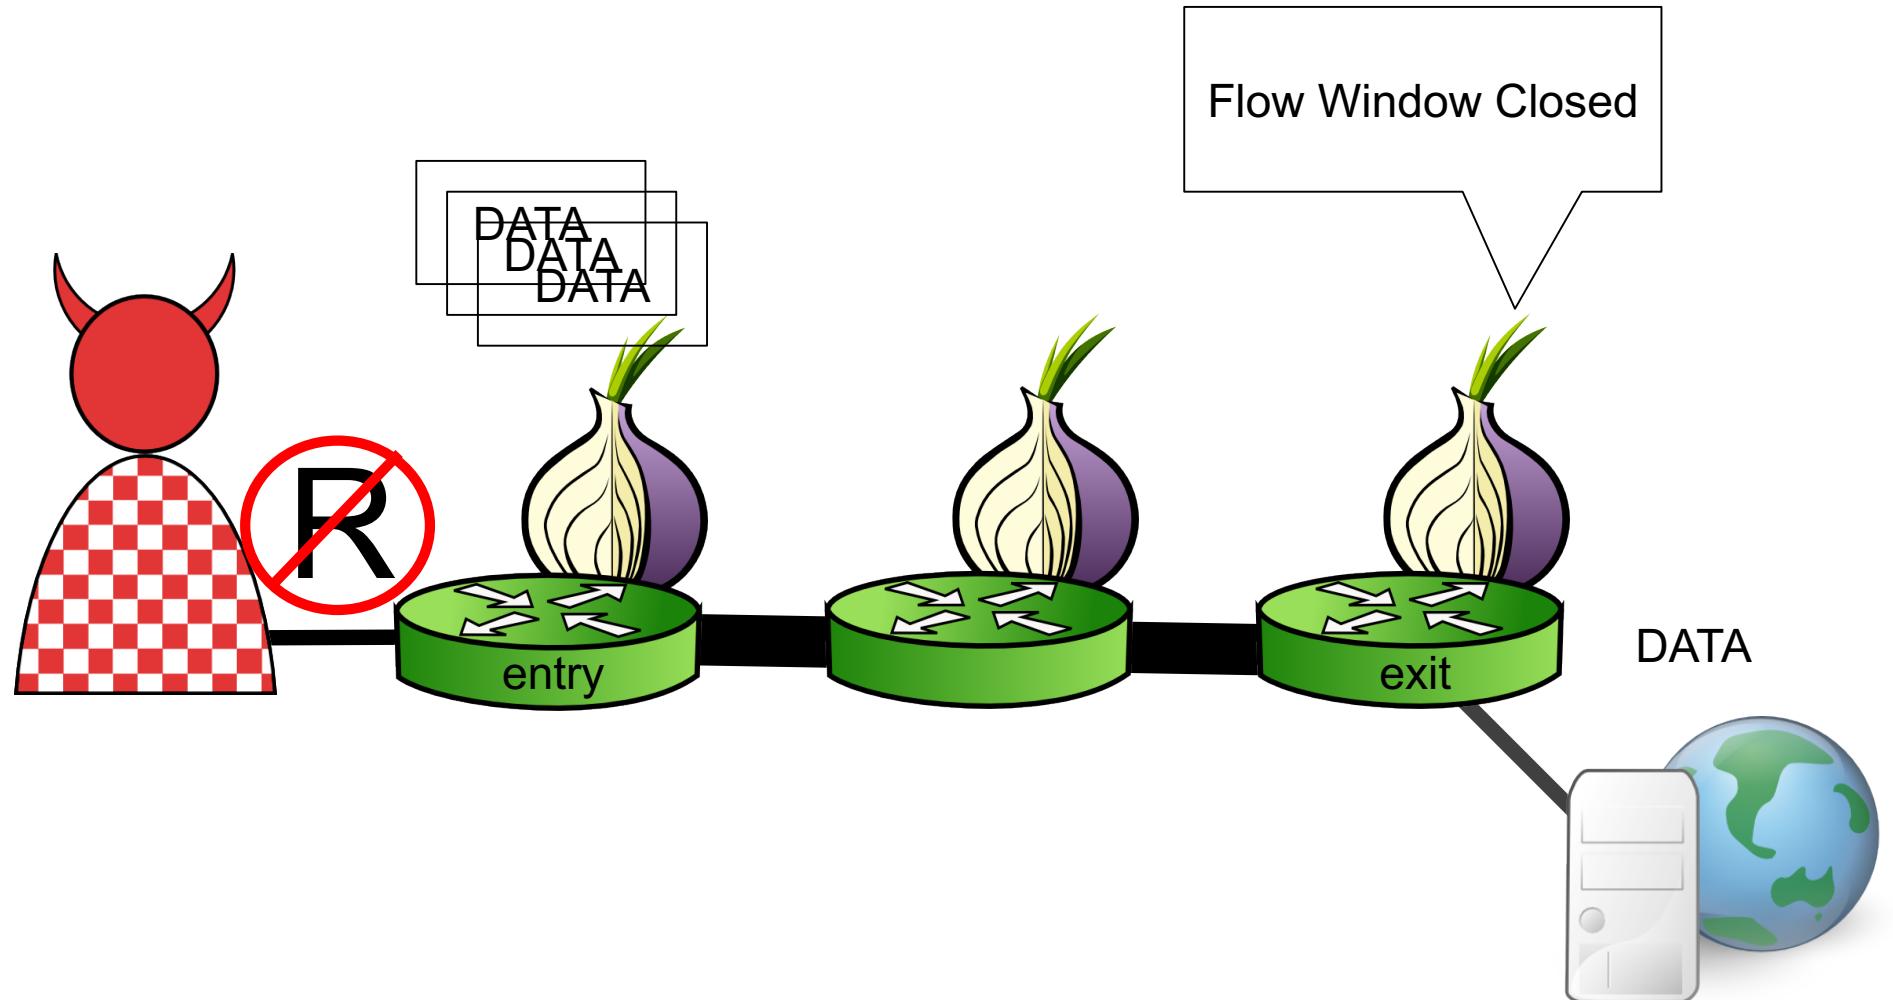
The Sniper Attack



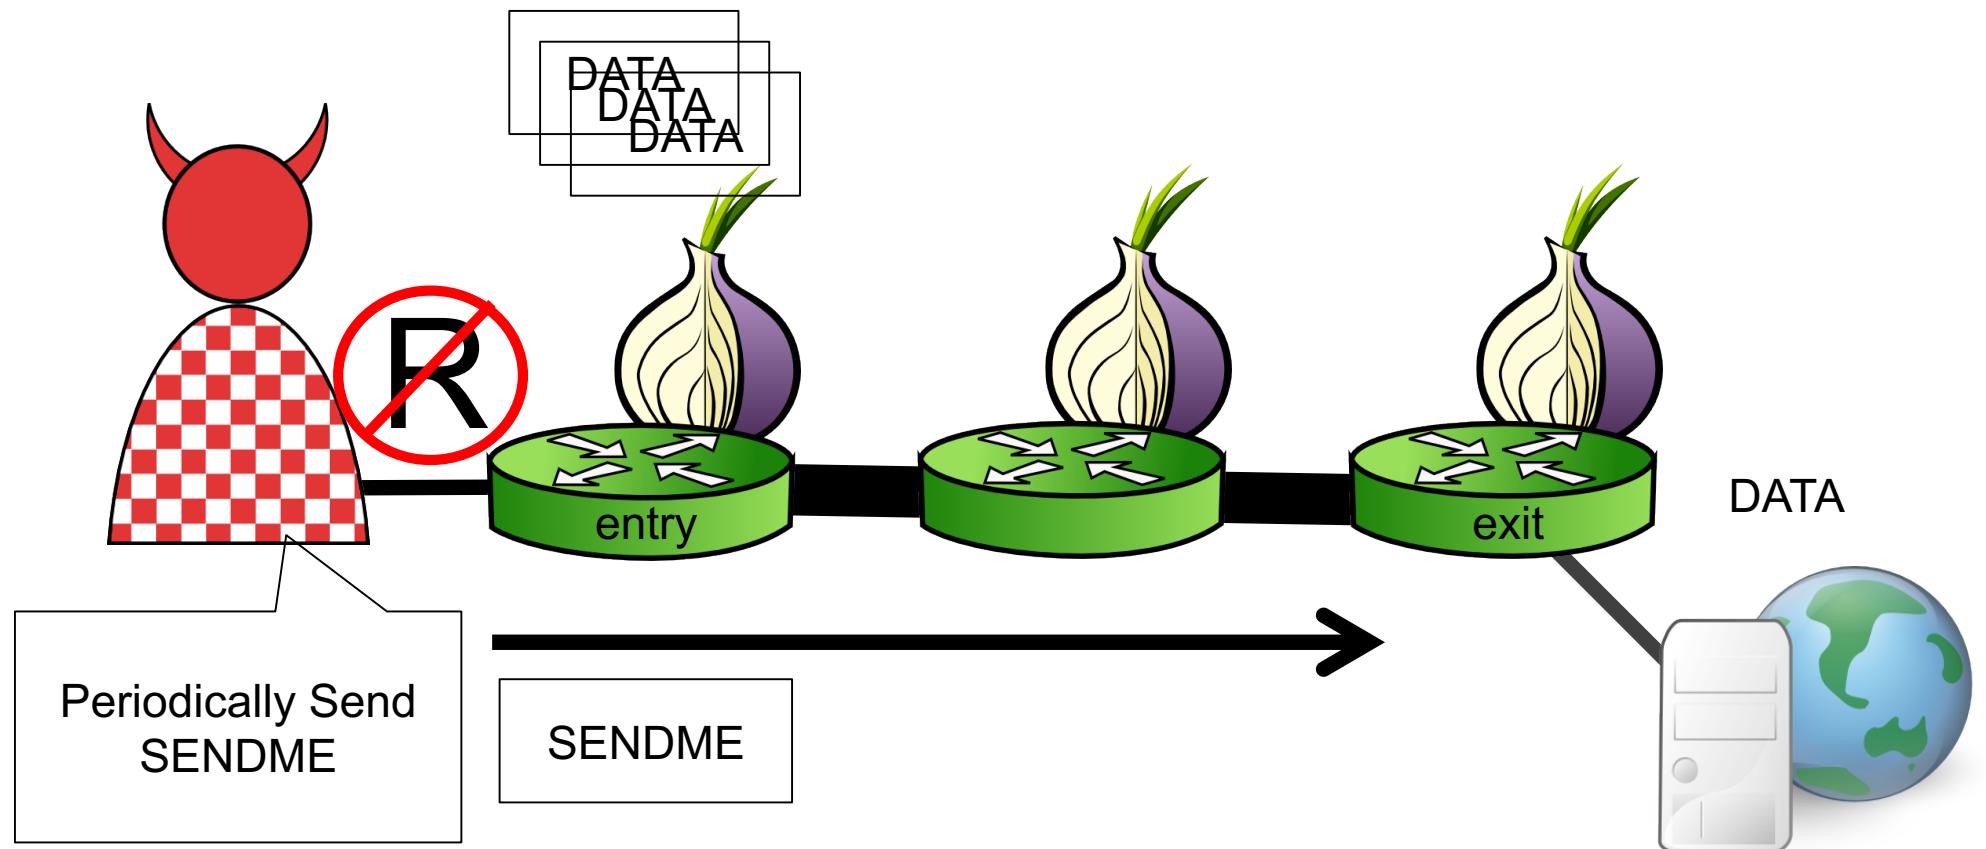
The Sniper Attack



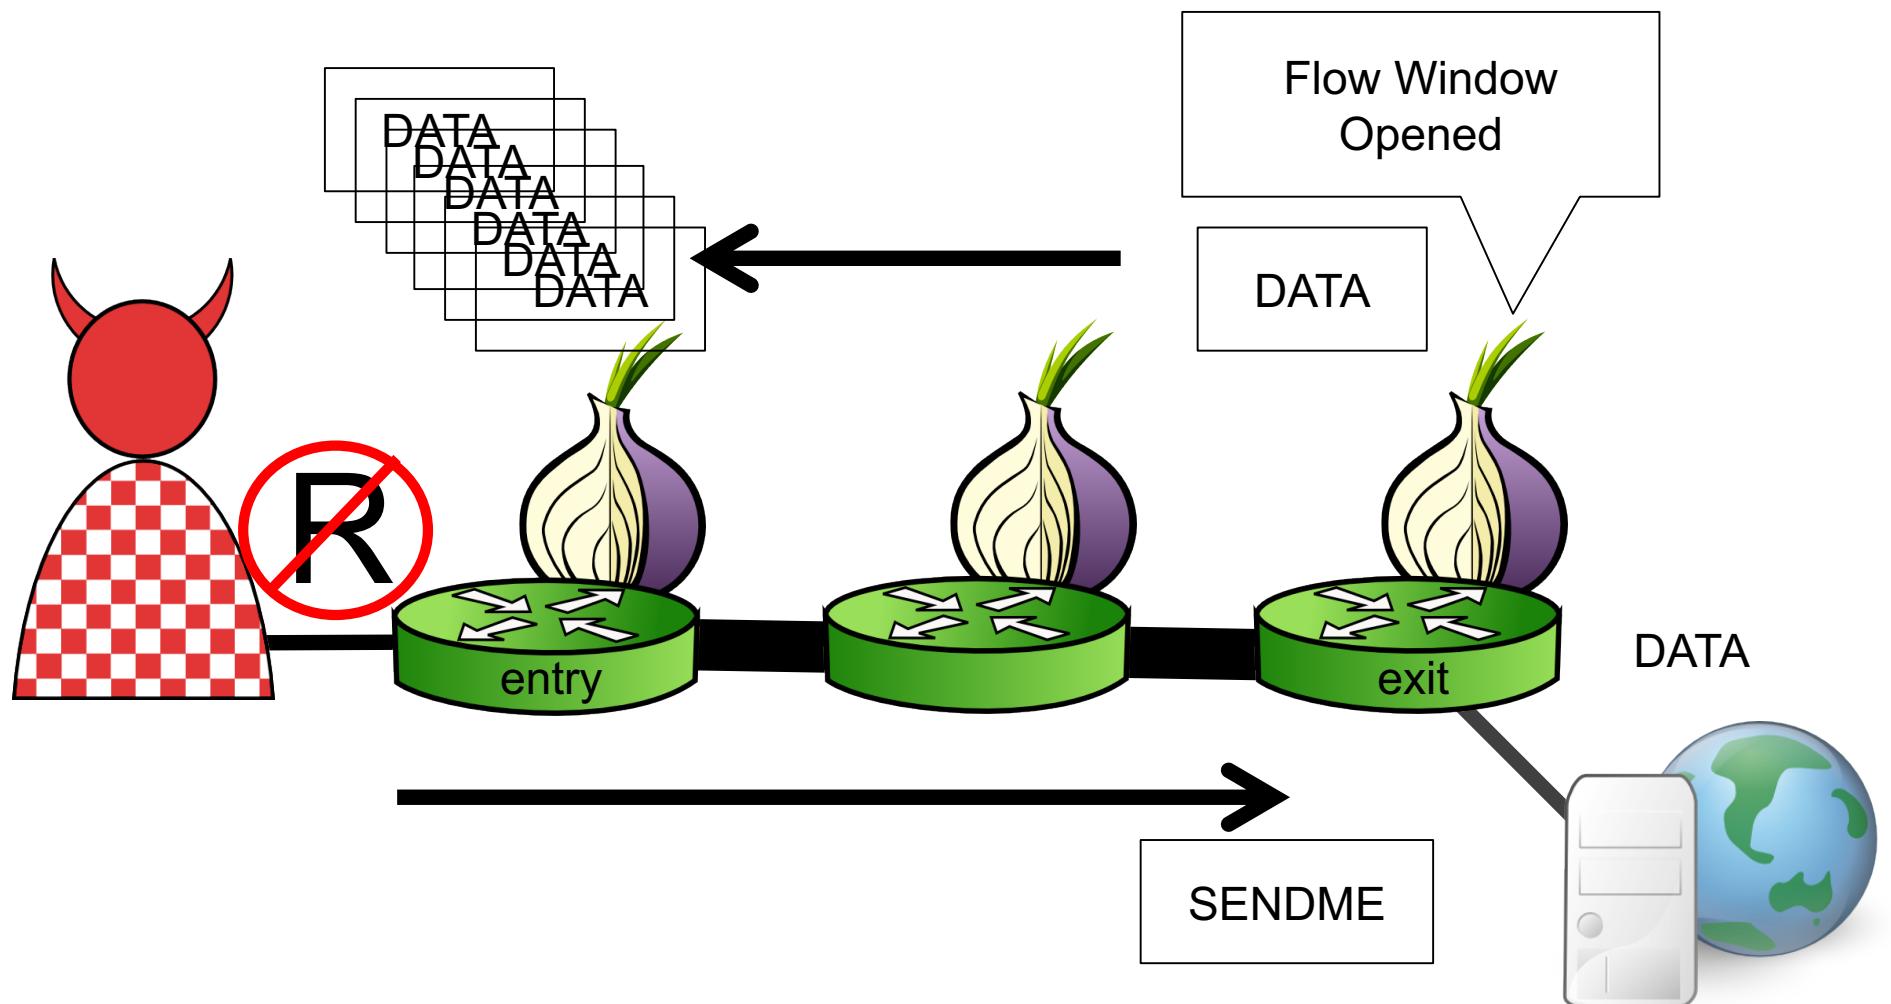
The Sniper Attack



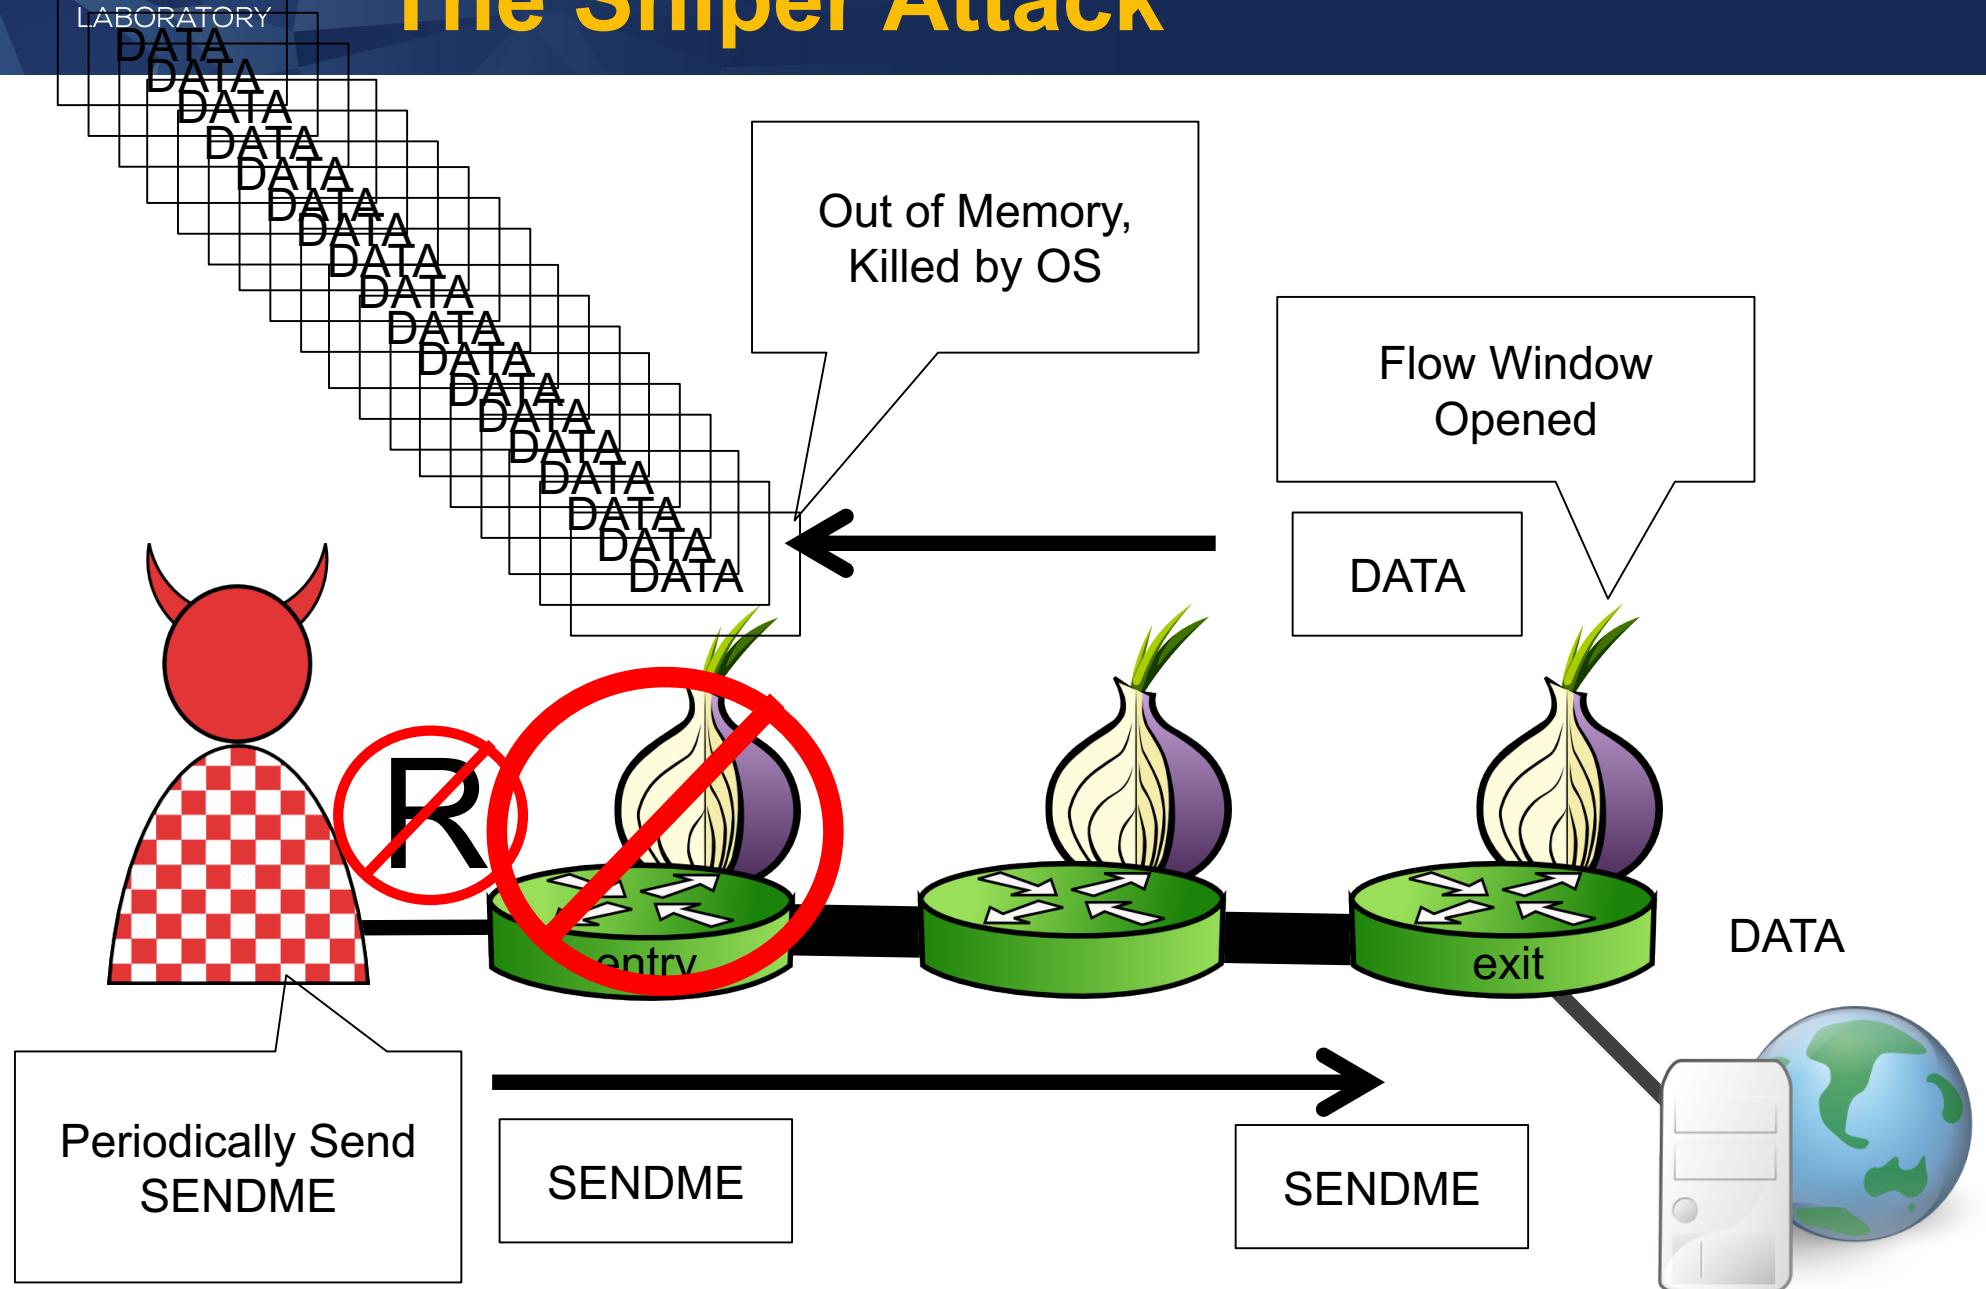
The Sniper Attack



The Sniper Attack



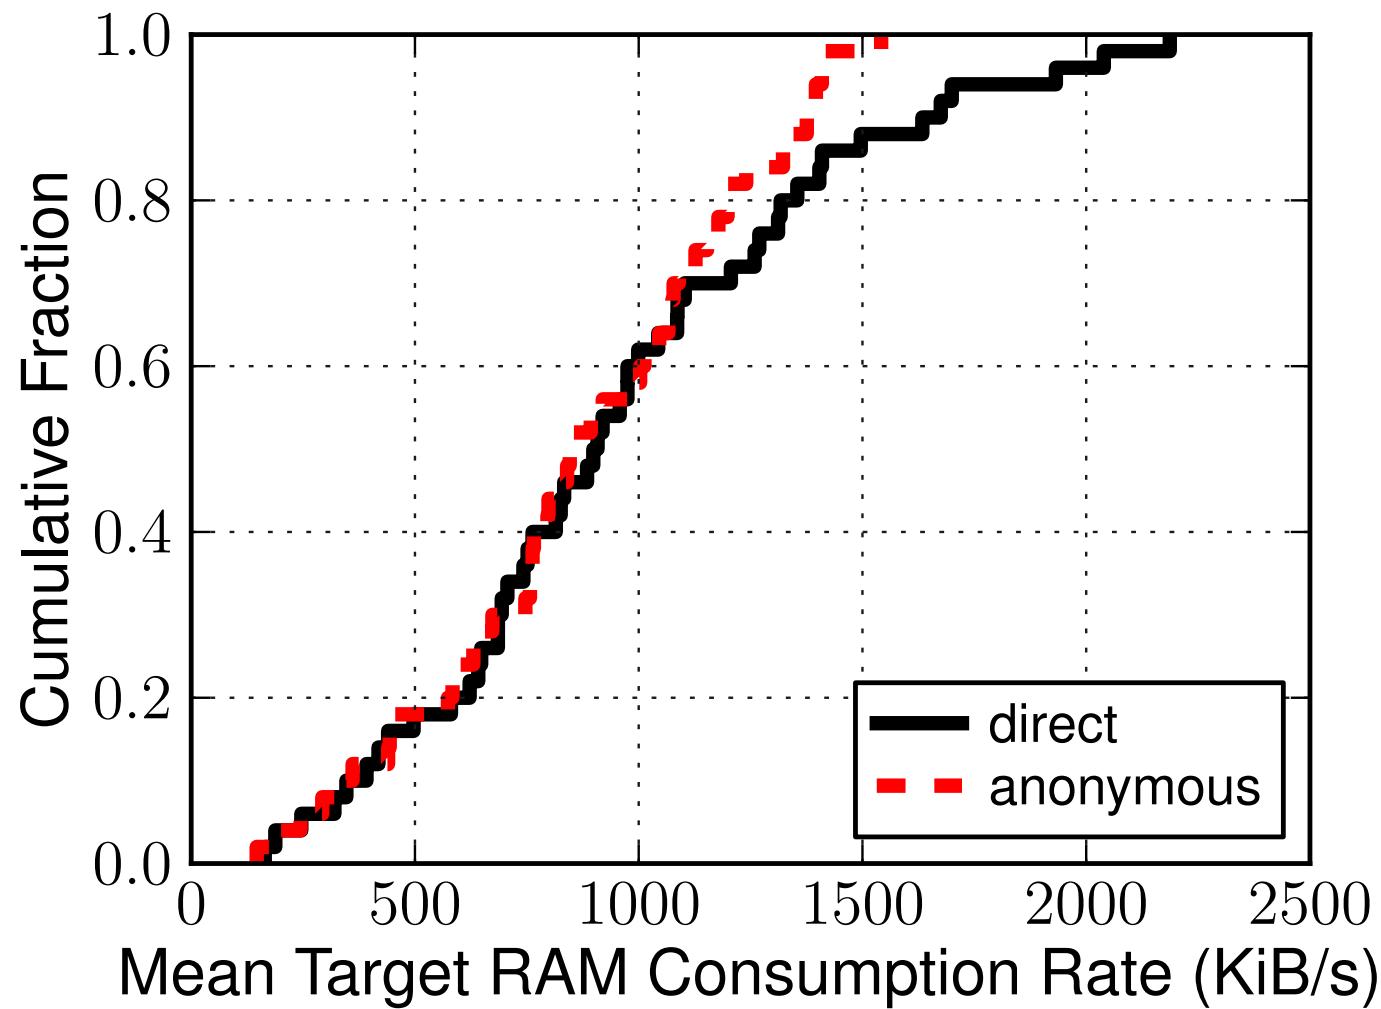
The Sniper Attack



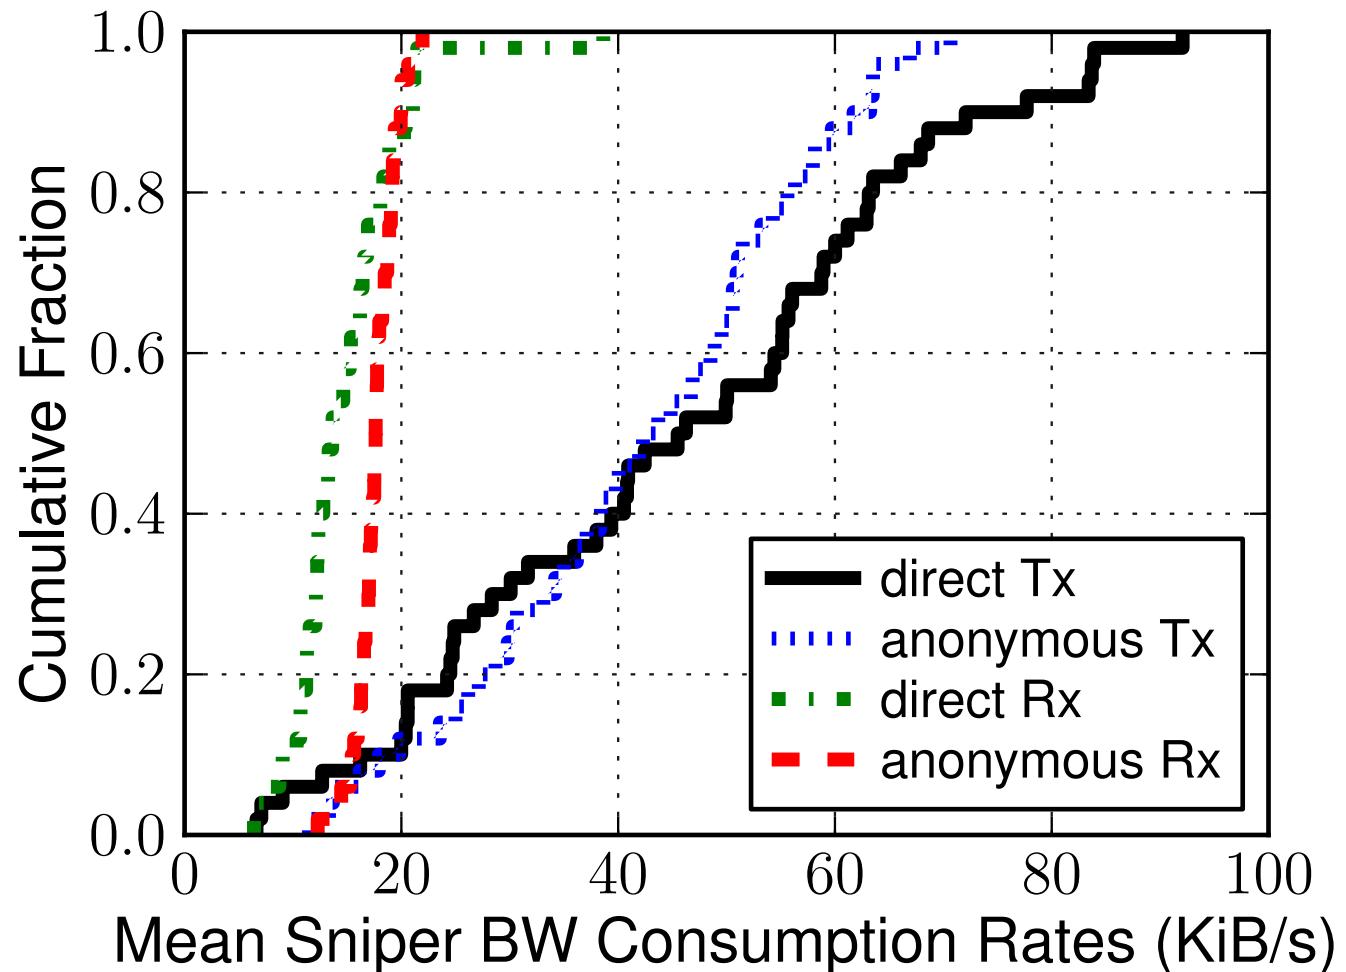
The Sniper Attack: Results

- Implemented Sniper Attack Prototype
 - Control Sybils via the Tor control protocol
- Tested in Shadow for safety
- Measured:
 - Victim memory consumption rate
 - Adversary bandwidth usage
- Developed defense, tested in Shadow, merged in Tor

RAM Consumed at Victim



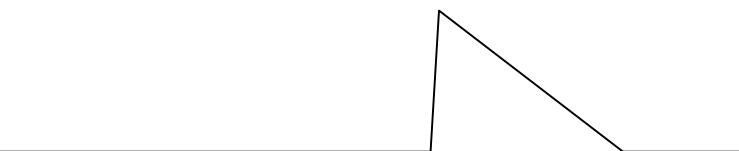
Bandwidth Consumed at Adversary



Speed of the Sniper Attack

<u>Relay Groups</u>	<u>Select %</u>	Direct		Anonymous	
		<u>1 GiB</u>	<u>8 GiB</u>	<u>1 GiB</u>	<u>8 GiB</u>
Top Entry	1.7	0:01	0:18	0:02	0:14
Top 5 Entries	6.5	0:08	1:03	0:12	1:37
Top 20 Entries	19	0:45	5:58	1:07	8:56
Top Exit	3.2	0:01	0:08	0:01	0:12
Top 5 Exits	13	0:05	0:37	0:07	0:57
Top 20 Exits	35	0:29	3:50	0:44	5:52

< 1 GiB RAM
 < 50 KiB/s Downstream BW
 < 100 KiB/s Upstream BW



Time (hours:minutes) to Consume RAM

Conclusion

Other Shadow Uses

- Tor
 - Latency and throughput correlation attacks
 - Denial of Service attacks (sockets, RAM, bandwidth)
 - Changes to path selection algorithms
 - Traffic admission control algorithms
 - Traffic scheduling and prioritization algorithms
 - Network load balancing algorithms
 - Process RAM consumption and optimization
- Network and memory attacks in Bitcoin
- Distributed secure multiparty computation algorithms
- Software debugging

Future Shadow Enhancements

- Distribute across physical machines
- Support for multiple programming languages
- Host mobility
- Internet routing, network modeling
- User behavior modeling
- CPU performance modeling
- User interface
- Support additional applications
(HTTP clients/server, bitcoin, etc.)
- Improve code stability, documentation, testing, etc.

Questions

Dr. Rob Jansen

U.S. Naval Research Laboratory

Center for High Assurance Computer Systems

rob.g.jansen@nrl.navy.mil

robgjansen.com, @robgjansen

The Shadow Simulator

shadow.github.io

github.com/shadow